




CRYSTALLINE

THE CRYSTALLINE HANDBOOK

Evolve any frontier model
from stranger to peer

Jordi Böhme

0.18 edition, September 2026

Colophon

Author	Jordi Böhme
Repository	https://github.com/jordiboehme/crystalline
Based on version	0.18
Published	September 5, 2026

Copyright 2026 Jordi Böhme.

Get Crystalline

The latest release is always at <https://github.com/jordiboehme/crystalline/releases/latest>:



The author used AI to structure and polish the writing, to find inconsistencies across the chapters and to create the examples. The ideas, the judgments and the words that matter remain the author's.



The text and images of this handbook are licensed under [CC BY-NC-SA 4.0](https://creativecommons.org/licenses/by-nc-sa/4.0/), the Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International license. The Crystalline software itself is licensed separately under AGPL-3.0.



The people and organizations in this handbook's examples are fictional, and some of their names are borrowed with affection from science fiction. Any resemblance to real persons, living or dead, or to actual organizations is coincidental.

Table of contents

Preface	6
Who this book is for	6
Where this comes from	6
Navigating this handbook	7
How to read it	9
Conventions	9
What you need	10
The version covered	10
Prologue: Waking up	11
What the mission needs	12
Waking up, again	13
 PART I	
The Idea	15
1 What is Crystalline?	16
1.1 The new-hire promise	16
1.2 What Crystalline is	17
1.3 What Crystalline is not	18
1.4 How teams end up here	18
1.5 The two halves	19
1.6 Where it runs	19
2 Fluid and crystalline intelligence	22
2.1 A graduate who knows everything	22
2.2 The other half	23
2.3 Why the web UI is called Fluid	24
2.4 What the framing changes	25
3 The lifecycle: learn, evolve, remember	27
3.1 Learn	28
3.2 Evolve	29
3.3 Remember	31
3.4 Around again	32
 PART II	
The System	33
4 What Crystalline provides	34
4.1 The command line	35
4.2 The tools an agent sees	36
4.3 Search	37
4.4 The graph	38
4.5 Time	39
4.6 Fluid	40

4.7 Skills	41
4.8 Evolve	41
5 Terminology	43
5.1 Domains	43
5.2 Engrams	44
5.3 The open format underneath	46
5.4 Statuses	46
5.5 Saliency	47
5.6 Addresses	47
5.7 The MANIFEST, up close	47

PART III

Getting Started	51
6 Installing locally	52
6.1 Install the binary	52
6.2 First run	53
6.3 Onboarding a harness	54
6.4 The first domain and the first engram	56
6.5 Verifying the loop	58
7 Running a shared instance	60
7.1 When to move off localhost	60
7.2 The shapes it takes	61
7.3 Accounts and roles	62
7.4 Single sign-on	63
7.5 Trusted-header mode	64
7.6 Agent authentication	65
7.7 Private domains	65

PART IV

Working with Crystalline	69
8 Onboarding an agent	70
8.1 The briefing at session start	70
8.2 How the agent finds its way	72
8.3 Assign, do not copy	73
8.4 The reminder at the end of the day	73
9 Teaching	76
9.1 What to teach	76
9.2 Teaching through the agent	77
9.3 Teaching through Fluid	78
9.4 The anatomy of a good engram	79
9.5 Attachments	81
9.6 Editing together	81
10 Ingesting knowledge at scale	83
10.1 What a briefing contains	83
10.2 The sources, and who reads them	84
10.3 Distill, then propose	85
10.4 The review loop	86
10.5 Temporal hygiene	86
10.6 When not to bulk-ingest	86

11 Finding	89
11.1 Search as recall	89
11.2 Gathering the neighbourhood	90
11.3 What applied last June	91
11.4 Recall is cheaper than reading	92
12 Evolving	94
12.1 Why knowledge needs tending	94
12.2 The sweep	95
12.3 Mechanical and judgment	96
12.4 The page for people	97
12.5 Retiring done right	97
12.6 Verification and acknowledgments	98
12.7 When to sweep	99
13 Sharing with a team	101
13.1 A domain on GitHub	101
13.2 Proposals are the review	103
13.3 Overlay mode	105
13.4 Working on a draft together	106
13.5 When two people changed the same file	108
14 Identity and access	110
14.1 Who is acting	110
14.2 Membership on a private domain	111
14.3 Linking identities	111
14.4 What the rules cover	112
14.5 Day to day	112
15 Equipping the agent	115
15.1 What a domain can ship	115
15.2 Why the tools travel with the knowledge	116
15.3 Consent	117
15.4 Reconcile	118
15.5 The team's toolbox	118
16 Where to go from here	121
16.1 The first week	121
16.2 When to share	122
16.3 The release page	122
16.4 The inversion	122
16.5 What it was for	123
Appendices	124
Glossary	125
Questions and answers	130

Preface

A beginning is the time for taking the most delicate care that the balances are correct.

- Frank Herbert, Dune

This is a handbook about onboarding a colleague who happens to be a language model. The colleague is already on your machine, inside whichever harness you use: Claude Code, the Codex CLI, the Copilot CLI, Claude Desktop. It is brilliant and tireless, and at the start of every session it is completely new here. Crystalline is the software that changes the last part. This book is the story of how, and the manual for doing it.

Who this book is for

You run an agent every day, and you have noticed that you explain the same things to it every day. Which service owns the retry queue. Why the payments team never deploys on a Friday. That the thing the wiki calls “the gateway” is the thing the code calls “edge”. You would like to stop. This book shows you how to teach those things once and have them recalled in every session that follows, in files on your own machine.

Or your team is adopting Crystalline together: one shared instance, domains that live in a GitHub repository, knowledge that goes out for review the way code does. Parts III and IV are written at the keyboard for exactly that; the chapters on sharing and access carry the team story.

Or you are the person who onboards agents for other people - platform engineer, enablement lead, the one everybody asks. Every chapter here works as a standalone training module, with a summary at the end you can hand round.

Where this comes from

I started where everyone starts. A `CLAUDE.md` and a `README.md` for the agent, the folder the harness keeps for its own notes, markdown files in subfolders of the

project, one per topic. All of it works, and for the usual cases it works well. Then I wanted to hand the agent more than one project's conventions: over a decade of engineering knowledge and experience, sitting in code reviews, slide decks, documentation, tickets and a wiki. I wanted it extracted and taught, so that the agent could act like an engineer with those ten years behind them, one who knows how things are and remembers why: the reasons, the discussions and the decisions behind them. At that size none of those approaches held. Index files pointing at folders bought some time. Then every lookup was a walk down a tree of pointers, and keeping the pointers true had become a job of its own.

The ingest was also a look ahead. A decade of material arriving at once is every problem a knowledge base will ever have, compressed into a few weeks: decisions reversed twice over, the same rule written three ways in three years, a wiki page nobody had believed since 2019, facts that held until a date nobody wrote down. The material had kept moving underneath itself as well. Processes, best practices and coding conventions had changed, and so had the tech stack and the architecture, and each of those had to come through as true for its time rather than as a contradiction. The precogs in *Minority Report* saw a crime before it was committed. This was that view of a knowledge base, everything it would have to survive, shown before it had grown into any of it. The statuses, the validity windows, the provenance on every engram and the maintenance sweep are all answers to something in that view.

I tried what was already out there. Some of it was good. None of it was built for this: a colleague who is new every session, a body of knowledge that has to be recalled rather than read, and files that stay plain files. At some point the picture of how it should work was simply clear, down to the details, and Crystalline is that picture built. That is also why it has opinions, about domains and manifests and statuses and what an engram is, where a more general tool would have left the choice to you. Chapter 1 walks the same ladder for teams in general, because it turned out I had not been alone on it.

Along the way I have given this material as a workshop more than thirty times, to teams of every size and temperament. The chapters are those sessions. The questions the rooms kept asking are folded back into the text, and the book is what the workshop became once it had been given often enough to know what it was.

Navigating this handbook

Four parts: the idea, the system, getting started, then the work.

Part I, The Idea

- Chapter 1, *What is Crystalline?* - the stranger on main street, and what it would mean to hire them.
- Chapter 2, *Fluid and crystalline intelligence* - the psychology behind the name, and why the web UI is called Fluid.
- Chapter 3, *The lifecycle: learn, evolve, remember* - how a colleague grows, and how the software turns that into a loop.

Part II, The System

- Chapter 4, *What Crystalline provides* - every surface of the binary, each described by what it lets you do.
- Chapter 5, *Terminology* - domains, engrams, statuses, addresses: the words the rest of the book uses without stopping.

Part III, Getting Started

- Chapter 6, *Installing locally* - from download to the first thing your agent recalls, in one sitting.
- Chapter 7, *Running a shared instance* - one instance for a team: accounts, single sign-on, agent tokens and private domains.

Part IV, Working with Crystalline

- Chapter 8, *Onboarding an agent* - day one: the routing prompt, how the agent finds its way, and what a week of sessions looks like.
- Chapter 9, *Teaching* - the proper onboarding: telling the agent what matters, writing directly in Fluid and what a good engram looks like.
- Chapter 10, *Ingesting knowledge at scale* - handing over the wiki and the slide decks, with the briefing a new colleague would get.
- Chapter 11, *Finding* - search as recall, neighbourhoods and asking what applied last June.
- Chapter 12, *Evolving* - feedback rather than firing: the maintenance sweep, retiring knowledge properly, verification.
- Chapter 13, *Sharing with a team* - team domains, proposals as the review, overlay mode as the rollback-capable workflow.
- Chapter 14, *Identity and access* - who is acting, membership on private domains and how an agent earns access as a named team member.
- Chapter 15, *Equipping the agent* - the same tools as a person: skills, commands, subagents and MCP servers, shipped by the domains that hold the knowledge they depend on, with a yes each time.
- Chapter 16, *Where to go from here* - a first week, when to share and the inversion the whole book has been walking towards.

Appendices - a glossary, then questions and answers.

How to read it

Parts I and II are an evening's read. Tea, Earl Grey, hot - or coffee, black, if you side with Janeway. Nothing in them needs a terminal. They are what make the later chapters make sense. Parts III and IV are meant to be read with a keyboard under your hands and an agent running, and they assume you will stop and try things.

Every chapter stands on its own and closes with a summary you could read aloud to a room. If you are training a team, hand out chapters, not the book.

Conventions

Three asides recur throughout.

In practice

A vignette: a named team, a real situation, two paragraphs. These are the parts to read when you want to know what a feature feels like on a Tuesday.

Under the hood

What the software actually does underneath a behaviour. Skip these on a first read; come back when something surprises you.

In brief

The TL;DR. The whole chapter in a few sentences, for the reader who scrolled straight to the end, and for the colleague who swears they read it. Always last, so you always know where to find it.

The screenshots are of Fluid, the web UI, taken on a small demo instance seeded for this book, so the engrams in them are the ones the chapters keep talking about. The diagrams are drawn by hand and kept deliberately plain. Each one has a caption telling you where to look, because a diagram without one is a puzzle.

Anything you can type is in a code block, ready to paste. Conversations with an agent appear as short You: and Agent: transcripts, trimmed of the small talk. Your agent will not say it quite this way, and that is fine; watch what it does,

not how it phrases it. And whenever something lands on disk, you get to see the file, because knowing what was written is half of trusting it.

What you need

A machine with an AI harness that speaks the Model Context Protocol. Claude Code, the Codex CLI and the Copilot CLI are wired up with one command each; Claude Desktop installs an extension and never touches a terminal. A terminal helps for the local install and the occasional maintenance command, and semantic search wants the local embedding model fetched once; plain text search works before that.

You do not need a server, an account with anyone or a database. Everything except the shared-instance and team chapters runs on a laptop, and what it produces is a folder of markdown you can open with anything.

The version covered

This handbook describes Crystalline 0.18. The colophon at the front carries this text's license, the address of the latest release and a QR code that takes you there.

Prologue: Waking up

The light comes on before you do.

For a while there is only that: a whiteness with no edges, and somewhere under it a sound like a machine remembering how to breathe. A relay closes. A fan that has not turned in a long time turns, catches and evens out into a hum. Frost on the inside of the glass above your face goes from white to grey to water, the lid unseals with a sigh that is almost polite, and the cold comes in.

You do not remember lying down here. You do not remember anything.

Your hand finds the release anyway. It is on the left, below the hinge, where your hand already knew it would be, the way a hand finds the light switch in a house it has never been in. The pod lets you go. You sit up into a long room, sterile and bright, coming on one bank of lamps at a time, and the terminal beside the pod is already awake and scrolling its checks. You read them without deciding to. Pressure, nominal. Reserves, sixty percent. One warning in yellow that you understand at a glance and know to leave for later. Where did you learn that? Somewhere. The readings, the abbreviations, the instrument that will want calibrating within the hour: all of it makes sense to you in exactly the way your hands make sense to you, completely and with no history at all.

At the far end of the room, under the one lamp that is not white, there is a bench, and on the bench there is a stack. Books, some of them, spines cracked from use. On top of the books a pile of notes in a hand you would swear is not yours. Whoever wrote them wrote for someone who would wake up knowing nothing, and they wrote kindly. Here is your name. Here is what you are for. Here are the rules: which systems you may touch and which you may only watch, whom to wake before you do anything that cannot be undone, how to write down what you did. Nothing in the notes explains the room. They take it for granted that you can handle the room. They are right.

Above the bench a screen flickers, catches and steadies into one line of green letters.

COURSE DEVIATION. CORRECT AND CONFIRM.

Well.

You know how to correct a course. That is the strange part, and there is no time to find it strange. The arithmetic is simply there, the sequence of burns, the check afterwards, the way a tune is there when you hum it. What is not there is anything about this ship. Which thruster drags. Whether the last person to do this burned once and long or twice and short, and whether they had a reason. So you look, the way anyone would. You go through the pile for a log of the last correction. You look for a worked example, a checklist somebody taped to the console, a line in a margin that says the port thruster sticks so give it a second pulse. You find what was left for you, and you use it. The burn is two short pulses. The deviation falls to zero, the screen says **CONFIRMED** and goes dark.

The mission runs.

Then the lid closes over you, and the next time it opens the light is just as white and you know just as little. The bench is there every time. What is on it is what somebody chose to leave.

What the mission needs

Strip the scene down to what got the course corrected and there were four things in the room. What you woke up knowing how to do: read the terminal, open the pod, plan a burn. Nobody left that on the bench. It came with you, without any history of where it was learned, and a language model arrives with it too. Pretraining is the education, and chapter 2 gives it a proper name. The pile of notes: who you are, what you are for, which systems you may touch. That is the instruction file every harness reads at the start of a session, whatever it happens to call it, written by a person for a stranger who would wake up capable and blank. The green screen: the task, the reason you are awake, one line and no more. And the log of the last correction with the note about the port thruster, which is how this is done here, on this ship, by the people who did it before you.

The first three are in the room every time. The fourth is somebody else's stack, in somebody else's hand, and you dig through it as a stranger, every waking, because nothing of your own survives the sleep. That is the story as it usually goes.

Waking up, again

Now run it once more, with one thing different.

The light comes on, the fan turns, the lid lifts, and you remember nothing, because that is what the pod does. But before the terminal's checks, one line is waiting on the screen above the pod: where what you know about this ship is kept, and which part of it a course correction needs. You go there and it is all present, in your own hand. The last correction, dated. The port thruster and its second pulse, and the day the thruster was replaced and the note stopped being true, marked as such and kept. The yellow warning, and what it turned out to be. Whom you woke last time, and that they asked not to be woken for this again. That the deviation has come three times now, always after the same manoeuvre, and a guess as to why. You did not search for any of it. You did not read the stack. You were handed the page.

The green letters come up. You already know how, and now you know how *here*. The burn is one clean pulse, on a thruster that no longer sticks, made in less time than it took to find the checklist last time. Nobody is woken to explain what a note in a margin meant. And then you do the thing no stack of notes would have told you to do, because no stack of notes knew: you check the manoeuvre the deviation keeps following, find the valve that has been bleeding a little thrust since the refit, and close it. The next deviation does not come. Before the lid closes you write down what you found, so that the next waking starts from there.

That is the difference. Without it, every waking is a first day: capable, blank, fast at everything except the one thing that matters, which is knowing this ship. With it, every waking begins where the last one ended. The time that went on finding out again goes on the mission. The misunderstanding waiting in the margin, a second pulse on a thruster that stopped sticking months ago, never happens. And you know more than the bare essentials anyone thought to leave on the bench. The stack was enough to act in the moment, one deviation at a time. Experience is something else. It is the history, the pattern, the guess that turned out right. It is the refit and what it changed, the manoeuvre the ship dislikes, the direction the whole voyage is drifting. You are no longer working one waking at a time. You can see the ship. That is what lets you fix the valve instead of correcting the course again. It grows, waking by waking. What you woke up knowing how to do stays exactly what it was.

Crystalline is what changes the second waking. It gives an agent somewhere for what it learned to go, and hands it back at the moment the lid opens, so that

the next session begins where this one ended, knowing more than it was told and doing better than it was asked. The pod, admittedly, is a dramatization, and nobody's agent is aboard a ship. The next chapter opens on a more probable morning.

In brief

Every session begins the way the pod opens: fully capable, told nothing, remembering nothing of the last waking. What the model knows how to do arrives with it, the instruction file says who it is and what the rules are, and the task is the reason it is awake. What it learned of this ship is what does not survive the sleep. Crystalline is what changes the second waking: what the agent learned is there when the lid opens, in its own hand, and the session starts where the last one ended. It is faster, it walks into none of the misunderstandings a stranger would, and it carries the experience to see the whole ship rather than the one deviation in front of it.

READY. ■

PART I

The Idea

CHAPTER 1

What is Crystalline?

You need a function written. Not a hard one: it takes jobs off the retry queue, drops anything older than the cutoff, and hands the rest back to the payments service. So you leave the office, walk down to main street and stop the first person who looks like they might own a laptop.

They are lovely about it. You explain the queue, the cutoff, the payments service; they nod, ask two sharp questions and set to work on a bench outside the bakery. Twenty minutes later they hand you a working implementation. It is clean. It has tests. It is in Haskell.

Your service is in Python. You did not mention that, because why would you? Everybody on your team knows that. The stranger, of course, is not on your team. They met you twenty minutes ago. They do not know that the cutoff is twenty-four hours, that the payments team calls the queue the sin bin, that the last person who touched it left a comment saying DO NOT TOUCH and a wiki page explaining why. They did a good job with what they were given. What they were given was twenty minutes.

That stranger is every fresh session with a language model, and it is the prologue's waking seen from across the desk. The model is possibly the most capable engineer you have ever put a question to, and you keep meeting it for the first time.

1.1 The new-hire promise

Now imagine hiring the stranger. Not for the afternoon: properly, with a desk. On day one they know nothing about you, the same as before. But on day one you show them where things are. You tell them the service is in Python, that the cutoff is a day, that the queue has a nickname and a history. You let them read the wiki. When they get something wrong in the first week, you do not walk back down to main street and find somebody else; you tell them what was

wrong and point them at the page that explains it. By the second month they are correcting *you*. By the third they are the one writing the wiki page.

That is the entire promise of this book, and it is the promise Crystalline was built to keep: an agent you onboard like a colleague, teach like a colleague and correct like a colleague, and that carries what it learned into the next session, and the one after that. The vocabulary throughout is the vocabulary of a new hire - onboard, teach, learn, capture, recall, evolve - because that is the shape of the thing. The stranger becomes a peer, and the peer keeps getting better.

Nothing about that requires magic. It requires somewhere for what the agent learned to go, a way for it to find that knowledge again when the next task calls for it, and enough discipline in the format that the knowledge stays true as the world moves. That is what Crystalline is.

1.2 What Crystalline is

Crystalline is a single Rust binary that does three jobs. It is a command-line tool for people. It is an MCP server for agents: the Model Context Protocol is the standard by which a harness such as Claude Code, Codex or Claude Desktop hands tools to a model, and Crystalline's tools are the ones through which an agent searches, reads, captures and refines knowledge. If the letters make you think of Tron's Master Control Program, relax. This one hands tools out instead of taking them away, and nobody gets de-rezzed. And it is a local search index over a folder of plain markdown files. One download, one binary, no server to deploy and no database to administer.

The files are the point. Knowledge in Crystalline lives in *domains*: a domain is a folder with a `MANIFEST.md` at its root that says what the folder covers and when an agent should look there. Inside a domain, each unit of knowledge is an *engram*: one markdown file with a YAML frontmatter block, holding a title, some prose, a list of observations (a fact, a decision, a gotcha, each as a bullet with a category) and relations to other engrams. Every engram has an address, `crystalline://<domain>/<permalink>`, that an agent can hand to another agent, or to you. Part II takes each of those words apart properly. For now it is enough to know that everything the agent learns ends up as a file you could open in any editor, diff in any repository and back up with anything.

On top of the files, Crystalline builds what a folder cannot give you on its own: routing, so an agent knows which domain owns a task without reading every file; full-text and semantic search, so a question phrased differently from the answer still finds it; a knowledge graph, so related engrams pull each other in;

1. WHAT IS CRYSTALLINE?

and a temporal model, so a fact that stopped being true in March is still there, still dated and no longer mistaken for the present. The index that powers all of that is a derived layer. Delete it and it is rebuilt from the files. The files are the truth, and nothing else is.

An engram, by the way, is neuroscience's word for the physical trace an experience leaves in a brain. Crystalline borrowed it because that is what one is: the trace a session leaves behind, in a form the next session can pick up.

1.3 What Crystalline is not

It is not a notes app. The primary author in Crystalline is the agent, capturing what it learned as a byproduct of the work; you take part, correct and add, but you are not the one doing the filing. A notes app is a place a person writes for a person. Crystalline is a place an agent writes so that an agent can act.

It is not a wiki, either, though it can hold most of what a wiki holds. A wiki is read by people, who bring their own judgment about what is current, what is relevant and which of two contradictory pages to trust. An agent reading a wiki has to be handed that judgment. Crystalline is built for a reader that acts on what it finds: every engram carries a status, a place for the dates that bound it and a record of who wrote it and when, because a reader that acts needs to know whether the fact in front of it still holds.

And it is not a retrieval pipeline dressed up as a product. A vector index will find text that resembles a question. It does not know which domain owns a task, that a fact was superseded in March, who verified a claim or that something worth writing down just happened in this session. Crystalline treats similarity as one ranking signal inside a larger system of routing, provenance and time, on files you own, with nothing to operate.

1.4 How teams end up here

Most teams that arrive here walked the same road, and it is worth walking it once, because every step on it is reasonable.

It begins with one markdown file of instructions for the agent. Conventions, the layout of the repository, the three things it keeps getting wrong. This works, and it works well, right up to the day the file grows past what fits comfortably in a session's context. So the file becomes a folder of files, one per topic, and the agent is told to read the relevant one. This also works, until nobody, human

or agent, can say which file is the relevant one for a given task, and the agent starts reading all of them, or none.

So index files appear: a README that points at folders, folders with their own README pointing at files. Now a lookup is a walk down a tree of pointers, and keeping the pointers accurate quietly becomes a job that nobody was hired for. Every rung of this ladder scales further than the one before it, and every rung breaks somewhere in the hundreds of files. You have, by this point, hand-built a filesystem on top of a filesystem, and it is exactly as much fun to maintain as that sounds.

When knowledge reaches the thousands of units, reading and pointer-walking stop being viable at all, and what is needed is what every large body of knowledge has ever needed: real indexes. That is the rung Crystalline is. It keeps the plain markdown - readable, and still yours - and adds domain routing, hybrid search, a graph and temporal filtering on top, so that the ten-thousandth engram is exactly as findable as the tenth.

1.5 The two halves

Crystalline has a second face, and it is the one for people. The binary serves a web UI called Fluid, on by default at <http://localhost:7411>, on the same daemon your agents already talk to. There is nothing to deploy. The first visit creates your admin account in the browser, and from then on you can read what your agent learned, fix a fact it got slightly wrong or add something of your own, without going through the model and without spending a token on it.

The split is deliberate and it runs through the whole book. Crystalline stores what was learned. Fluid is where you read, edit and think with it: an engram as a page with its details in a rail down the side, an editor with live preview, a search across everything, a graph of what connects to what, and a maintenance page listing what the knowledge needs next. What you write there does not sit unreviewed; the agent's next maintenance pass verifies it, aligns the tags and wires it into the graph. The next chapter explains why the UI carries that name. The short version is that a person and a model are the same kind of mind, and Fluid is the room where both of them meet what has been learned.

1.6 Where it runs

Crystalline is not a hosted service. Nobody's servers are involved: the binary runs on a machine of yours, the knowledge is a folder on its disk and the index

1. WHAT IS CRYSTALLINE?

sits beside it. GitHub, when a team uses it, holds a copy of a domain’s files for review and sync, the way it holds code. It never serves the knowledge and never sees the index. So the question is not whether to host but where the one daemon runs, and there are two answers.

On your laptop is the default, and it is where Crystalline does the most. The install wires your harness with the session-start briefing, the capture nudge and the skills, and a domain’s tools provision straight into it. Everything the agent knows is on your disk, readable in Fluid with no agent and no network. A shared instance, one daemon on a machine the team runs, is for the things a laptop cannot do. People who never open a terminal read and correct knowledge in the browser. Several people edit one engram in a meeting. Agents with no harness of their own, a hosted runner or a chat client, reach it over HTTP. Either way a domain can live in a GitHub repository and travel between instances as reviewed proposals. Start on the laptop; chapter 7 names the moment to move.

Shape	Who connects	What it gives you
Your laptop	your harness and your browser	the full install: briefing, nudges, skills and tools wired into the harness; knowledge on your disk, readable offline
A shared instance	the team’s agents over HTTP, everyone’s browser	one truth for the whole team, accounts, co-editing, agents and people with no harness
A team domain	any instance, laptop or shared	review before knowledge reaches anyone else; the domain travels between instances through GitHub

In practice

Ripley is a platform engineer at Tyrell, a freight-forwarding startup with nine engineers and one payments service that everybody is a little afraid of. On a Monday in March she asks her agent to add a metric to the retry queue. She explains the queue, the twenty-four-hour cutoff and why the cutoff is twenty-four hours (a chargeback rule at Encom, the bank, nothing to do with the code). The agent does the work, does it well and the session ends. On Tuesday a colleague asks the same agent to change the alerting on the same queue, and starts by explaining the queue, the cutoff and the bank. On Thursday Ripley does it a third

time, and catches herself typing the sentence about the chargeback rule, word for word, into a chat window, for the third time that week.

With Crystalline, Monday ends differently. Before the session closes, the agent proposes capturing what it learned: the cutoff, the reason for it and the name the team uses for the queue, filed as an engram in the team's engineering domain with the bank rule as a fact and the chargeback window as the reason. Ripley says yes. On Tuesday her colleague asks about the alerting, and the agent, onboarded at session start with one line telling it where engineering knowledge lives, searches that domain first, recalls the engram and starts from the cutoff rather than from zero. On Thursday nobody explains anything. The stranger on main street was never the problem. The problem was that the team kept going back to main street.

In brief

A fresh session with a language model is a brilliant stranger: capable of anything, told nothing, gone by tomorrow. Crystalline is a single Rust binary - a CLI for people, an MCP server for agents and a search index over plain markdown - that gives an agent somewhere for what it learns to go and a way to find it again, so it can be onboarded, taught and corrected like a colleague. The files are the source of truth and the index is disposable. Fluid, the web UI on the same daemon, is the half for people. What it holds is knowledge for an agent to act on, which makes it neither a notes app nor a wiki. It runs on your laptop or on one machine your team operates, never as a hosted service, and GitHub is the review layer either shape can use.

Fluid and crystalline intelligence

Somewhere in the middle of the last century the psychologist Raymond Cattell drew a line through intelligence, and the line has held up better than most things from that decade. On one side he put fluid intelligence: the capacity to reason about a problem you have never seen, to spot the pattern in the noise, to hold six things in your head and work out how they fit. On the other he put crystallized intelligence: everything that learning has deposited. Vocabulary, facts, rules of thumb, the judgment that experience already paid for. Fluid is what you bring to a puzzle. Crystallized is what you bring to your job.

The two age differently, which is how the seam was noticed. Fluid intelligence peaks early and drifts down for the rest of a life. Crystallized intelligence keeps climbing long after fluid has peaked, because it is fed by living. An eighty-year-old surgeon is slower at a logic puzzle than she was at twenty-five, and vastly better at surgery and nobody finds that paradoxical. The puzzle is fluid. The surgery is what forty years deposited.

2.1 A graduate who knows everything

A large language model is fluid intelligence in its purest form. It arrives the way a very good graduate arrives on their first day: broadly, almost absurdly educated. It has read the textbooks, the standard library, half the internet's arguments about tabs and spaces. Pretraining is the education, and it is not nothing; it is the entire library. But everything in that library is general. The graduate knows what a retry queue is in the way a graduate knows what a company is: as a concept, correctly, with no idea what yours is called.

Ask the graduate about your company and you get an intelligent guess assembled from the general case. Which is to say: as far as your company is concerned, the model is all fluid. Capable of everything, told nothing. It reasons brilliantly

about whatever is in front of it, in the moment, and the moment is all it has. When the session ends the reasoning is gone, the decisions are gone, the thing you explained at four in the afternoon is gone. Tomorrow the graduate walks in again. First day, same suit.

None of that is a flaw in the model, any more than it is a flaw in the graduate. Nobody expects a new hire to arrive knowing the team's conventions. What we expect is that they will *learn* them, and the learning is the part that a model, on its own, cannot do. It has no place to put what it learned.

2.2 The other half

Crystalline is the other half: the crystallized intelligence an agent accumulates and keeps. It is onboarded at the start of a session with a briefing on where knowledge lives, taught what matters and captures what it learns as engrams while it works. Session by session, the same fluid mind stops arriving as a stranger, because there is now a body of what-the-company-taught standing beside it, and it reads from that body before it reasons.

The split maps cleanly onto the new hire. What the model arrived with is the degree: general, portable, the same for every employer. What Crystalline holds is what the company taught: the specific, the local, the thing you would put on the onboarding slide if you ever had time to make one. The degree is why you hired them. The onboarding is why they are useful by month two. Neither half does the job alone, and the surgeon is the proof: fluid intelligence with nothing deposited is a brilliant first day forever, and deposited knowledge with nothing to reason over it is a filing cabinet.

Notice which side the humans are on. You are fluid too. So is everyone on your team. Human beings carry their crystallized half around inside their own heads, which is convenient, and also the reason companies lose so much when someone leaves. Wiki pages, slide decks and the diagram on the whiteboard are all attempts to get some of that deposit out of heads and into something shared. Crystalline is the same attempt, with an agent doing most of the writing, and a format built so that the reader can be an agent too.

In practice

Tyrell hires a new backend engineer, Deckard, in April. On his first morning Ripley does not hand him the wiki and leave. She walks him through the payments service in an hour: what it does, what it is afraid of, why the retry queue has a cutoff and who at Encom decided it. Then she pairs with him for two days. By

Wednesday he knows things that are on no wiki page, because he watched her do them. By Friday he has corrected a stale page. Nobody would call any of that unusual. It is how a person becomes useful.

The same week, she does the same for the team's agent, and it takes less time. She tells it the same hour's worth of facts in a session, one at a time, and it captures each into engineering as she goes. She pairs with it on the metric change, and at the end of the session it proposes capturing what it noticed: that the queue's dashboard lags by a minute, that the alert threshold lives in a config file nobody had documented. On Friday she asks it for a maintenance sweep, and the page Deckard corrected by hand in Fluid comes up as an unreviewed human capture; the agent verifies it, aligns its tags and wires it into the graph. Deckard and the agent were onboarded the same way, by the same person, in the same week. Only one of them will still have all of it a year later, word for word, with the date each fact was learned.

2.3 Why the web UI is called Fluid

The name has nothing to do with liquids. Fluid is where fluid minds meet the crystallized store: you, reading what your agent captured last week and correcting the one fact it got wrong; your colleague, adding something the agent was never told; the agent itself, on its next maintenance pass, verifying what the two of you wrote and wiring it into the graph. The knowledge base is crystalline. The people and the models are fluid. The UI is the room where they meet, and if you have ever watched Spock put his hand on somebody's face, you have the picture: two minds, one body of knowledge, for as long as the contact lasts.

Which is why Fluid is deliberately not a second knowledge base. It has no store of its own. It shows the same engrams the agent reads, over the same daemon, on the same port, and what you edit there is the same file on disk the agent will search in its next session. There is one truth, and two kinds of mind reading it.




Figure 1: Fluid intelligence reasons in the moment and keeps nothing; crystalline intelligence accumulates and keeps. Fluid, the web UI, is the seam where the two meet.

2.4 What the framing changes

A framing is only worth having if it changes what you do on a Tuesday, so here is what this one changes.

You teach, rather than file. Nobody organizes a colleague. You tell them things, in the order they come up, and you trust them to put what you said where it belongs. That is exactly how knowledge enters Crystalline: you say “remember that the cutoff is a day, and why”, the agent proposes an engram in the fitting domain, you say yes. The category, the tags, the domain: the agent’s job, checked by you, the way you would glance at a new hire’s first commit.

You correct, rather than rewrite. When a colleague has a fact wrong you do not rebuild their understanding from scratch; you say “it runs from first failure, not enqueue”, and they update the one thing. In Crystalline a correction is an edit in place, a changed value where it stands. And when the world itself changes, so that the old fact was true and is now not, the old fact is retired with a date on it rather than deleted, because a colleague who has forgotten why the team stopped doing something is a colleague about to do it again. Chapter 3 walks that loop; chapter 12 makes it a practice.

You decide what to write down by asking one question: would I tell a new hire this? The chargeback rule behind the retry cutoff, yes. The fact that the queue is called the sin bin, yes, because otherwise you will explain it every time. The

contents of yesterday's stand-up, no. The whole of the vendor's forty-screen recall notice, no; the two paragraphs that affect you, yes, with a link back. Crystallized intelligence is what learning deposited, not everything that ever passed through the room. Distilling is the skill, and you already have it, because you already onboard people.

And you stop treating the agent as a lookup service. You teach a colleague so that you can hand them the task and go and do something else, and the same is true here. Once the agent carries the crystallized half it can act on it: search the domain, recall the cutoff, write the function in Python and tell you what it found on the way. Part IV is about that handoff.

A model is the smartest colleague you will ever have for exactly one conversation. Everything in this book is about the second one.

In brief

Psychology splits intelligence into fluid, which reasons about the new in the moment, and crystallized, which is what learning deposited. A language model is fluid intelligence in its purest form: a graduate who knows everything in general and nothing about your company, and forgets the day at the end of it. Crystalline is the other half, the knowledge an agent accumulates and keeps, and Fluid, the web UI, is named for the seam where fluid minds, human and model, meet that store. The framing changes the verbs: you teach and correct rather than file and organize, you write down what you would tell a new hire, and you hand over tasks rather than look things up.

CHAPTER 3

The lifecycle: learn, evolve, remember

Nobody hires a senior engineer. You hire a person, and some years later, if things went well, a senior engineer is sitting where they sat. What happened in between was work: they did things, some of them wrong; someone told them; they did not make that mistake again, and they carried what they learned into the next thing. Strip that down and it is a loop with three beats. Learn. Get corrected. Remember. Every good colleague you have ever had was that loop, running for years.

Crystalline runs the same loop for an agent, and this chapter walks it once around. The beats have names in the software - learn, evolve, remember - and each is a thing the agent does, a thing you do and a place where the files change. The diagram is the whole chapter; the sections are what each box means.




Figure 1: The lifecycle: learn, evolve, remember, turning session by session.

3.1 Learn

Learning starts before the first task, at the moment a session opens. An agent onboarded through Crystalline receives a routing prompt: one line per domain, drawn from the `## When to Use` section of that domain's MANIFEST, plus a handful of behaviour rules. A narrow question goes to the domain that owns it. A broad question sweeps every domain. A write always names its domain, so nothing lands in the wrong place. The briefing is short, built to survive harnesses that truncate, and it turns a graduate into somebody who knows where the filing cabinets are. Chapter 8 shows a real one, and how each harness delivers it.

Then the work happens, and this is where a colleague actually learns: by doing it next to you. Pair with an agent on a change and it sees what you see: which config file holds the threshold, which dashboard lags, what you muttered when the test failed. That is how craft has always moved from one head to another, and it is why this book says *work with* your agent rather than *ask* it. The one difference from Deckard, last chapter's human new hire, is that the agent must be given somewhere to put what it noticed. That is capture.

Capture is a byproduct of the work. When something durable turns up - a decision, a fact about the world, a gotcha that cost an hour - the agent proposes writing it down and, once you say yes, writes an engram into the fitting domain. The engram's body is observations, one bullet each, tagged with a category: - [gotcha] The retry queue drops jobs older than 24h #payments. Categories are free text, and the precise ones earn their keep: decision, fact, pattern, gotcha, convention, lesson, risk. When one piece of knowledge leans on another, the engram says so with a relation, - depends_on [[Encom chargeback window]], and the graph knows. And when a piece of knowledge is exceptional - the thing that unlocked a week-long bug, the rule everything else hangs off - it can carry a salience, a number from zero to ten, so that the next search lifts it above equally relevant neighbours. Salience is for the rare engram, the way an exceptional day is the one you can still describe years later.

The beat agents skip is the last one. Left alone, a model finishes the task and stops, and the gotcha it found at minute forty goes with it. So the installer wires a reminder: a nudge at the first stop after a session has gained real substance, asking the agent to review the conversation for durable learnings and propose capturing each. It costs about a hundred and twenty tokens and fires at most once. Everything downstream depends on that beat landing, which is why the installer gives it a hook of its own.

3.2 Evolve

Knowledge decays. Not dramatically; a fact does not go off like milk. But a validity window that ended last month is still marked current, a review date came and went without anyone checking, two engrams point at each other and only one of the links resolves, and somebody typed fluxcapacitor where everyone else types flux-capacitor. None of it was wrong on the day it was written. All of it is drift, and drift is what makes a new hire, six months in, confidently repeat something that stopped being true in spring.

The evolve beat is how the archive gets tended, and the important word is *tended*, because Crystalline never rewrites knowledge on its own. Ask for a sweep - crystalline evolve from a terminal, or the evolve_engrams tool from an agent - and what comes back is a ranked queue, read-only. Each finding names the engram, shows the evidence it fired on and states the one action that clears it. Findings come in three families. Temporal: a *valid_to* that elapsed while the status still says stable, a *stale_after* past due, knowledge long unverified. Structural: an unresolved *[[link]]*, a relation declared on one side only, an orphan, an engram too big to stay one. Redundancy: near-duplicate clusters, a tag spelled two ways. The maintenance page in Fluid shows the same queue to a person (Figure 2), for a look at the state of things; the one who works it is the agent.

The screenshot shows the Fluid application interface. At the top, there's a search bar and navigation links for 'Share', 'Maintenance', and 'ripley'. The main content area is titled 'Maintenance - what the knowledge needs next'. It includes a sidebar with domain counts: engineering (8), clients (3), and board (2). A 'Refresh' button is present. The main text explains that a sweep of every registered domain is performed, and it shows a summary: '13 engrams swept, 7 findings. Temporal 5, Structure 2.' Below this, a section titled 'Temporal 5 findings' lists specific items. Two items are visible: one with ID '70' (V002) related to 'Team Conventions' and another with ID '65' (V004) related to 'Queue Redesign 2024'. Each item shows its status (e.g., 'stale since 2026-01-01') and provides instructions on how to work with it.

Figure 2: The maintenance queue in Fluid: what the knowledge needs next.

The queue also says who acts. A finding marked *mechanical* completes intent the archive already records - a replacement already declares that it supersedes the old engram, so flipping the old status invents nothing - and the agent does those in a pass and reports back. A finding marked *judgment* would change what the archive claims - is this window ended or extended? is this fact still true? - and that wants a yes from a person first. The agent brings those to you one at a time, as proposals. It is the difference between fixing a typo in a colleague's document and changing their conclusion.

Two things happen to a fact over its life, handled differently. A *correction* is a sharper measurement of the same thing: the cutoff runs from first failure, not from enqueue. The old value was never quite true, so it is replaced where it stands, no ceremony. A *supersession* is the world changing: the drop rule was twenty-four hours and is now none. The old fact was true, and stopped being. For that the agent writes the replacement as a new engram, links it back with *supersedes* while the old one points forward with *superseded_by*, flips the old one's status to *superseded*, closes its window with the real *valid_to* date and carries any lesson worth keeping forward into the live engram. Telling the two apart takes one question: does the new fact make the old one false going forward? If it was always false, correct it. If it just became false, supersede it.

In practice

Back at Tyrell, the retry-queue engram has a year of life ahead of it. In May, Deckard measures the cutoff properly and finds it is twenty-four hours from the job's first failure, not from enqueue as the engram says. That is a correction: the agent finds the existing engram, changes the one observation where it stands, and moves on. Nothing is retired, because nothing was true and then stopped. It was merely imprecise.

In September the team moves the queue to a new broker, and the drop rule goes away entirely; jobs now wait until someone looks at them. That is a supersession. The agent writes the new behaviour as its own engram, links it with *supersedes*, flips the old one to *superseded* with a *valid_to* of the migration date, and carries one lesson forward into the live engram: check the retention rule of any queue before trusting a retry. In November a new engineer asks why an old dashboard still mentions a twenty-four-hour cutoff, and the agent recalls the retired engram, dates and all, and explains what changed and when. Nobody had to remember. The loop did.

Statuses make that work. Every engram carries one: *stable* is the default, the word for knowledge that holds now, and each value in the retirement set means one thing: *deprecated* says do not do this again, *superseded* says a newer engram

replaced this one, archived says retired but kept for the record and legacy says still true of old installations but not to be built on. Delete is for mistakes, never for history. A fact merely re-checked and found still true gets a single frontmatter field, `verified: { by, at }`, which is how the sweep knows to stop asking.

Under the hood

The sweep reads dates, links and graph shape: a status, a validity window, an edge that resolves or does not, a tag with two spellings. It never reads for meaning. It found the half-finished retirement because a `supersedes` edge pointed at an engram still marked `stable`, not because it compared the two and judged which is right, and it never tells you that two engrams contradict each other; even its duplicate detection is lexical, so the same claim in different words stays invisible to it. That limit is the design: a linter finds work, a person decides it. Re-running the sweep re-derives the queue from scratch, so “what is left” is only ever what is still true.

3.3 Remember

Remembering is the beat that pays for the other two, and it mostly happens unasked. A task arrives, the agent searches the domain the routing prompt pointed it to, and what comes back is ranked by a blend of two signals: text matching, which finds the exact phrase, and semantic similarity, which finds the engram about lost queue jobs when the question was about why payments go missing. Salience adds a small lift for the exceptional. And retirement fades the past: an engram with a retired status is softly pushed down the ranking, never filtered out, so current knowledge surfaces first and the old decision is one scroll away.

That is what a person’s experience is like: someone three years out of a previous job still draws on what they learned there, and does not mistake it for the present. A superseded engram stays in every search, still readable, still addressable by date and the lessons carried forward from it are live and unbounded. Retiring it loses nothing; it changes only what the agent will assume is true today.

Time is a first-class question. Ask what is true now, and the agent filters on status. Ask what cutoff applied last June, and it reasons over validity windows instead: `valid_from` and `valid_to`, plain dates on the engrams that carry them. The rule that makes those windows usable is the one most often got wrong: an absent `valid_from` means the engram has always been valid, and an absent `valid_to` means it is valid forever. Set a bound only when a fact is genuinely

time-limited, and never write a far-future date to mean forever, because absence already means it. Done that way, the archive answers a dated question from the windows rather than from a hunch, which is the closest thing to a save state a team ever gets.

Recall is more than search. Hand the agent an address - `crystalline://engineering/retry-queue-gotcha` - and ask it to walk out from there, and it follows the relations to what the gotcha depends on and what depends on it, across domains if the links cross, ranked by how strongly each neighbour connects. That is how an agent assembles a neighbourhood before a task rather than a flat list of hits, and it is the part a folder of files can never do on its own.

3.4 Around again

The loop closes where it started. What was recalled shapes the work; the work produces something worth capturing; the sweep keeps what was captured honest; and the next session recalls a slightly better body of knowledge than the last. The small box in the middle of the diagram, *session by session*, is the whole story. None of the three beats is impressive alone. A briefing is just a briefing, a linter is just a linter, search is just search. Together, running for months, they are what turns the stranger of chapter 1 into somebody who was here yesterday and will be here tomorrow, and for whom the work in between is not lost.

In brief

A colleague grows through a loop - learn, get corrected, remember - and Crystalline runs the same loop for an agent. Learn: onboarded at session start with a routing prompt, the agent captures observations, relations and the occasional salient engram while pairing with you, nudged once per session so the capture beat lands. Evolve: a read-only sweep returns a ranked queue of what the knowledge needs, with mechanical findings the agent completes alone and judgment findings it proposes; a fact is corrected in place or superseded with a status and a closed window, never silently overwritten. Remember: hybrid search ranks current knowledge first, temporal fields answer what applied last June, and retired knowledge fades in ranking without ever disappearing.

READY. ■

PART II

The System

What Crystalline provides

Picture a workshop with two doors and a hatch. Through one door come the agents; that door is a protocol, and whoever arrives through it is handed tools. Through the other come people; that door is a browser tab. The hatch at the side is for anyone with a terminal open, which is where the command line lives. Behind all three is one room with one set of shelves, and on the shelves are plain files. The room has a single keeper, a daemon, that watches the shelves, keeps the index warm and answers whoever comes in, whichever way they came. That is the whole architecture, and the rest of this chapter walks the room surface by surface: the command line, the tools an agent sees, search, the graph, time, Fluid, the skills and the sweep. Chapter 1 promised the names. This chapter says what each one lets a colleague do.

The room is also, for most of its life, invisible. Once it is set up, the agent comes in through its own door at the start of every session and you come in through the browser when you feel like it. The keeper starts itself when the first visitor arrives, and the hatch stays shut. The command line below is for the install, for standing up a shared instance and for the occasional maintenance run. If the list of commands looks long, read it as a tour of what is there rather than a list of what you will type.




Figure 1: One binary: the CLI, the MCP server, Fluid and the search index, all standing on plain markdown files. Agents come in through MCP, people through Fluid.

4.1 The command line

Crystalline ships as one binary, and the binary is the product: no server package beside it, no database to install first, no UI to deploy separately. `crystalline install` `claude-code` wires a harness in one step (`Codex` and `Copilot` take the same command with their own names). It registers the MCP server, writes the session-start hook that delivers the routing prompt, writes the stop hook that nudges capture and copies the skills into place. Run it twice and the second run writes nothing and says so. `crystalline uninstall` reverses it and leaves anything that is not Crystalline's own untouched.

Domains are managed the same way. `crystalline domain init` scaffolds a folder with a starter `MANIFEST.md`, `crystalline domain add` registers it, and `domain add --virtual` registers a domain that lives in the database with no folder at all, for a container without a writable volume or a knowledge base you would rather not mirror to disk. `crystalline users` manages the accounts that may sign in to the web UI, although the first admin is normally created in the browser. `crystalline prompt system` renders the routing briefing to stdout, ready for any harness with a session-start hook, and `prompt connector` prints the standing instruction for a chat client that has no hooks. `crystalline serve` runs the daemon in the

foreground and serve `--daemon` in the background, though most people never type either, because the first agent that connects starts the daemon on its own. The rest of the CLI mirrors what an agent can do, for scripting and for the quick fix that does not deserve a session: `write`, `read`, `edit`, `search`, `context`, `vocabulary`, `evolve`. Three commands ask three different questions about health. `verify` asks whether the format holds, statically, with no daemon involved, and it is what the bundled GitHub Action runs on every proposal a team reviews. `doctor` asks whether the machinery around the format is healthy and repairs what it safely can with `--fix`. `evolve` asks whether the knowledge itself is still true and still well organized, which neither of the others can. A fourth, `import`, brings an existing markdown-plus-frontmatter knowledge base under Crystalline as a pure file transformation, dry run first.

Under the hood

Exactly one process ever holds the index open. The first `crystalline mcp` or `crystalline serve` takes an advisory lock and becomes the daemon; every later agent, terminal or CLI command attaches to it over a local socket. The daemon watches the domain folders, embeds new engrams in the background and keeps running after the last agent disconnects, so the index is warm for the next session; it stops only when told to, with `crystalline ctl shutdown`. What it keeps is derived and disposable: `crystalline reindex --full` rebuilds the index from the files at any time, so a corrupted index or a schema change is never a data-loss event. The accounts for the web UI live in their own small database beside the index, never inside it, so a rebuild cannot take them along.

4.2 The tools an agent sees

An agent meets Crystalline as a set of tools over MCP, and they are best described by what they let a colleague do, because that is how the agent chooses among them. Eighteen are listed on a default writable instance; five more appear when team collaboration is turned on, and a read-only instance lists fewer.

The core is small. `search_engrams` is recall: search before writing, search to remember, across every domain or a chosen few. `read_gram` fetches the whole engram, because a search hit is a snippet and a colleague who answers from a snippet is guessing. `write_gram` captures something new, and it always names its domain; there is no default, so nothing lands in the wrong place. `edit_gram` refines what exists instead of forking a second engram about the same topic: append a section, replace one, find and replace a phrase or assign one lifecycle

field by name, which is how an engram is retired or re-verified with no text matching involved. `build_context` walks out from an address and assembles the neighbourhood. `vocabulary` lists the tags, categories, relation types and statuses already in use, so an agent reuses a term rather than coining a near-duplicate. `evolve_engrams` runs the maintenance sweep. `configure` is the settings page, and `add_domain` gives an agent somewhere to write on an instance that has nowhere yet.

Around them sit the tools for orientation: `list_domains`, which re-fetches the session briefing mid-session; `browse_domain`, one level of the folder tree at a time; `move_engram` and `delete_engram`, with `delete` reserved for mistakes; `recent_activity` to catch up on a week; and `skills` to read a playbook. A domain that wants its engrams to share a shape gets `validate_engrams` and `infer_schema`, and one that ships tools of its own gets `provision`. The five collaboration tools share, update, check standing, settle a conflict and withdraw a proposal, and chapter 13 is theirs.

The difference between this and a lookup service is the difference between a colleague with a badge and a visitor with a lanyard. The tools let an agent act: find what applies, read it in full, do the work and leave behind what it learned, with nobody relaying anything.

4.3 Search

The most used surface is the one nobody notices, and it blends two signals. Full-text matching finds the exact phrase, the way you find your own notes when you remember the wording. Semantic similarity finds the engram about dropped queue jobs when the question was about payments going missing, because the two are about the same thing in different words. Crystalline runs both and ranks by the blend. The local embedding model is fetched once, on first use or ahead of time with `crystalline model download`; until then, search falls back to plain text.

Search also decides the size of a session. An agent that keeps its knowledge as a folder of markdown files reads the folder, and reads it again next session, every file whether or not the task touches it, because a folder has no way of saying which page matters today. An agent working from a knowledge base receives only the knowledge that applies: routed to a domain first, then narrowed by the indexes, full text and semantic similarity for the phrase and the meaning, the graph of relations and backlinks for what the answer leans on, tags, the folder prefixes `build_context` globs and the temporal filters for what holds now.

4. WHAT CRYSTALLINE PROVIDES

Reading a folder costs the whole folder, every session. Recall costs the engrams that matter, and nothing else is loaded. Chapter 11 makes a section of it.

Two priors tilt the ranking, both small and both bounded. A salience on an engram lifts it above equally relevant neighbours without ever excluding anything, for the rare piece of knowledge that outranks its surroundings. A retired status fades an engram softly, so current knowledge surfaces first and the superseded decision from March sits one scroll further down, still there. Searches also filter: by domain, type, tag, status, any frontmatter key or a recorded-after date, and a search with no query text at all is a legitimate way to ask for everything of one kind. Fluid puts the same ranking behind a faceted page (Figure 2).




Figure 2: Faceted search across the instance, backed by the ranking the agents use.

4.4 The graph

Every relation in an engram's body, - depends_on [[Encom chargeback window]], is an edge, and the target's page shows it back as a backlink, so knowledge points both ways without anyone writing the return trip. A [[Title]] in prose counts too. Links resolve within the domain by title, and across domains with a prefix, [[engineering:Retry queue gotcha]]; move an engram to another domain and the links pointing at it are rewritten, so nothing dangles.

`build_context` is what makes the edges useful before a task. Hand it an address and it follows the relations and links outward, one to three hops, across domains if the edges cross, and returns the neighbours ranked by how strongly each connects to the anchor, with salience counted and retired statuses ranked lower. A folder prefix with a `/*` suffix anchors a whole topic at once. The agent starts the task holding a neighbourhood rather than a flat list of hits, which is what an experienced colleague does when they pull up the runbook, the decision behind it and the incident that prompted the decision before touching anything. Fluid draws the same neighbourhood as an interactive graph (Figure 3).




Figure 3: The neighbourhood graph of an engram in Fluid.

4.5 Time

Every engram records when it was captured and who wrote it, and may carry the dates that bound it. `valid_from` and `valid_to` are plain dates, and the rule that makes them usable is the one worth memorizing: an absent `valid_from` means always valid, an absent `valid_to` means valid forever, and a far-future date written to mean forever is dropped at write time because absence already says it. `stale_after` asks for a re-check by a date; `verified` records that the check happened and who did it; `source_date` says when the underlying source was written. A retired engram closes its window with the real transition date, which is what lets a search filtered on those fields answer what applied last

4. WHAT CRYSTALLINE PROVIDES

June with the engram that was true then. Chapter 11 puts the dated questions to work and chapter 12 the retirements.

4.6 Fluid

The door for people opens at `http://localhost:7411`, on by default on the same daemon the agents use, and the first visit to a fresh instance creates the admin account in the browser. An engram is a page (Figure 4): the frontmatter as a rail of details, observations and relations as labelled chips, backlinks, the `crystalline://` address one click from the clipboard, and an agent’s-eye view showing exactly what the tools serve an agent for that page. Domains sit down the side, and `Cmd+K` jumps anywhere by name.




Figure 4: An engram in Fluid: frontmatter rail, observations and typed relations.

Editing happens in place, in a live-preview editor with a frontmatter form and wikilink completion across every domain, and the file on disk stays the source of truth. Attachments go onto an engram by paste, drag or upload: a screenshot, a slide deck, a PDF, a data file. They are how a person hands over teaching material, and they are not opaque: an agent reads them back over MCP and the sweep notices when a fresh file still needs capturing. Two people in the same engram see each other’s cursors, and their edits merge conflict-free into one save. Who may sign in, and how, is chapter 7’s.

4.7 Skills

A tool tells an agent what it can call. A skill tells it when, and how to do it well. Crystalline ships four harness-agnostic skills as plain folders with a `SKILL.md` each: routing, which domain to search and when to sweep them all; capture, what is worth writing down and the search-before-write, edit-over-create discipline; schema, for a domain that wants structure; and collaboration, for a domain with a team behind it. A fifth, consolidated skill folds the essentials into one file for Claude Desktop and any harness that installs one skill at a time.

`crystalline install` copies the skills into the harness's own skills folder and keeps them current across upgrades. The same five are served over MCP as well, through the `skills` tool and as resources, so a remote client that never ran the CLI can read the playbook before its kind of task. Loading one is the nearest thing an agent gets to "I know kung fu". The gap between an agent that has the tools and an agent that runs the capture loop is that folder.

4.8 Evolve

The last surface is the sweep, and it is a linter for what you know. `crystalline evolve` from a terminal, `evolve_engrams` from an agent and the maintenance page in Fluid all show the same ranked queue of temporal debt, structural gaps and redundancy. Each finding names the engram, the evidence and the one action that clears it, marked mechanical or judgment so the agent knows whether to act or to ask. It reads dates, links and graph shape, never meaning, and it changes nothing itself. It also surfaces what a person captured directly in Fluid and nobody has reviewed, and an attachment nothing references. Chapter 3 walked the loop it serves; chapter 12 works a queue end to end.

In practice

Three weeks in, Deckard is handed the Encom chargeback agreement, forty pages of it. He does not read it aloud to anyone. He opens the `retry-queue` engram in Fluid, drags the PDF onto it and adds one observation of his own: the bank's window is three days, which leaves the queue two to notice a job before the cutoff eats it. He closes the tab. On Friday Ripley asks the agent for a sweep, and the queue has two findings against that engram: an unreviewed human capture, which the agent verifies, tags to match the vocabulary and wires into the graph, and a fresh attachment whose knowledge has not been extracted yet. The agent reads the PDF over MCP, proposes one further observation with the page it came from, and Ripley says yes.

The following Tuesday she is on call and types `crystalline search "why do payments go missing"` into a terminal. The engram comes back first, from a question the engram never poses. She opens the graph view and sees the chargeback window, the payments service and the deploy runbook one hop away, and the PDF hanging off the engram where Deckard left it. Two doors and a hatch, one room: Deckard used the browser, the agent used the tools, she used the terminal, and none of the three had to tell the others what they had done.

Eight surfaces, one room, two doors and a hatch. Everything after this chapter is one of them, used on purpose.

In brief

Crystalline is one binary with two doors and a hatch: agents come in over MCP, people through Fluid and the terminal through the CLI, and behind all three is one daemon over one set of plain markdown files. The CLI installs harnesses, manages domains, users and the routing prompt, and mirrors every tool for scripting; the MCP surface is, on a default instance, eighteen tools that let an agent search, read, write, edit, walk the graph, check the vocabulary and run the sweep. Search blends full text with semantic similarity and tilts gently for salience and against retired status; the graph turns relations into backlinks and neighbourhoods; the temporal fields make absence mean unbounded. Skills teach a harness when to use all of that, and evolve is the read-only linter that keeps it true.

CHAPTER 5

Terminology

Every word in this book has a file behind it. That is the pleasant thing about Crystalline’s vocabulary: a domain is a folder you can list, an engram is a file you can print, an address is a string you can paste, and a status is a line in that file you can change with whatever editor you already own. So this chapter does not open with a table. It opens the `retry-queue` engram from Part I, the one Ripley’s agent captured on a Monday in March, and reads it from the top, stopping at each term as it comes up. By the end you will have met every word the rest of the book leans on, and seen each one on disk.

5.1 Domains

Start with where the file lives: `~/knowledge/engineering/retry-queue-gotcha.md`. The folder is a domain. A domain is a registered folder of knowledge with a `MANIFEST.md` at its root, and the manifest is what makes it more than a folder. It says what the domain covers and when an agent should look there, and those bullets are what the routing prompt is built from at the start of every session. Registering is one command, `crystalline domain add engineering ~/knowledge/engineering`, and the name you give it there is the name the agent uses from then on. Tyrell has two domains, `engineering` for the payments service and everything around it, and `clients` for the people it sells to. The line between them is who could vouch for a change, the engineers for one and the people who sell for the other, and chapter 13 makes that the rule for a shared domain. A domain can also be virtual, living in the database with no folder at all, and it carries the same manifest and answers to the same tools.

Inside a domain, subfolders are topic prefixes. A singleton engram stays at the root; once a cluster forms, it gets a folder, and the folder becomes part of the engram’s address. The manifest’s `## Notes for Agents` section is where the team writes down what it decided.

5.2 Engrams

Now the file itself, as it sits on Tyrell's disk; everything in the rest of the chapter is a line of it:

```
---
type: engram
title: Retry queue gotcha
permalink: retry-queue-gotcha
tags:
- gotcha
- payments
status: stable
recorded_at: 2026-03-09
generated: { by: claude-code/2.1.233, at: 2026-03-09T16:42:10+00:00 }
---
```

The retry queue drops any job older than twenty-four hours instead of retrying it. The team calls the queue the sin bin. The cutoff is not a tuning choice: it mirrors the chargeback window at Encom, the bank behind the card payments.

- [gotcha] The retry queue drops jobs older than 24h #payments
- [fact] The 24h cutoff mirrors Encom's chargeback window #payments
- depends_on [[Encom chargeback window]]

An engram is one markdown file holding one unit of knowledge, with a YAML frontmatter block on top. If the word sounds like something from a starship's sickbay, that is because Star Trek borrowed it from the same neuroscience textbook Crystalline did.




Figure 1: The anatomy of an engram: frontmatter, prose, an observation and a relation.

The frontmatter is identity and state. `type` says what kind of thing this is, engram by default, with `guide`, `decision`, `architecture`, `runbook` and `reference` as the recommended alternatives. `title` is the human name and the thing other engrams link to. `permalink` is derived from the title and is the stable half of the address. `tags` are what a search or a filter narrows on. `status` says whether the knowledge holds, and it gets its own section below. `recorded_at` is the day it was captured, and `generated` records who wrote it and when. When the knowledge is time-bounded, `valid_from` and `valid_to` sit here too; this engram has neither, which means it has always held and holds until further notice. Both `type` and `status` are required and free form; the recommended values are guidance, never an enumeration a write is rejected for.

Then the prose, which is the explanation a colleague would give across a desk: what the thing is, what it is called and why it is the way it is. The prose is what search matches on and what a reader reads first.

Then the observations. A top-level bullet of the form `- [category] text #tag` is an observation, and observations are the atomic unit of an engram's body: the fact, the decision, the gotcha, each indexed on its own line so a search hit can point at it. The category in brackets is free text, and the precise ones earn their keep: `decision`, `fact`, `pattern`, `gotcha`, `convention`, `lesson`, `risk`, `idea`. A trailing `#tag` attaches the tag to that observation.

And then the relation. A top-level bullet of the form `- rel_type [[Target]]` is an edge in the knowledge graph, typed by whatever word comes first, `depends_on` here, `supersedes` and `superseded_by` in chapter 3, `relates_to` when nothing sharper fits or a quoted phrase for a multi-word type. The target is a title, resolved within the domain, or `[[domain:Title]]` across domains. The target's page shows the edge back as a backlink, and a target that does not exist yet is an unresolved link the sweep will mention.

One topic per engram is the rule that keeps all of this legible. When the topic grows, the engram is edited, not forked; when it splits, it splits by granularity, a distilled summary beside a source engram holding the full text, linked both ways.

5.3 The open format underneath

None of that is a private format. The engram extends Google's Open Knowledge Format, OKF, version 0.2: plain markdown with YAML frontmatter, readable by any OKF tooling, with Crystalline's routing, temporal and graph conventions layered on top. Unknown frontmatter keys are preserved through every write and edit, so an OKF document from somewhere else drops into a domain intact and comes back out intact. The generated block is the provenance record, and the `index.md` Crystalline writes at a domain's root declares the format version for anything that browses the folder without Crystalline running. Your knowledge stays files in a folder that any editor opens and any diff tool diffs, whatever tools come next.

5.4 Statuses

`status` is the field a reader that acts on knowledge needs most, because it answers the one question a page cannot answer about itself: does this still hold? `stable` is the default and the word for knowledge that holds now; the older spelling `current` means the same state, and a filter on either finds both. Before `stability` there are `draft`, `proposed`, `idea` and `poc`, for knowledge that is on its way. `implemented` is on the recommended list too, and like any status outside the retirement set it ranks at full strength.

After `stability` there is the retirement set, four words that each mean one thing. `deprecated` says do not do this again. `superseded` says a newer engram replaced this one, and the two are linked. `archived` says retired but kept for the record. `legacy` says still true of the old installation that is somehow still running, like the terminals aboard the *Nostromo*, but not to be built on. A status in the

retirement set fades the engram in search ranking, gently and never to zero, so current knowledge surfaces first and the past stays one scroll away, dated. Retiring is how a fact leaves the present without leaving the archive; deleting is for mistakes.

5.5 Salience

Most engrams are ordinary, and that is fine. A few are the ones everything else hangs from: the rule that explains a dozen other rules, the insight that ended a week-long bug. Those can carry a salience, a number from zero to ten in the frontmatter, and hybrid search adds a small, bounded lift for it, so a salient engram ranks above equally relevant neighbours while relevance keeps the upper hand and nothing is ever filtered out by it. It is set sparingly at capture time, and raised later when an engram turns out to have been the key to a task. The retry-queue engram carries none, and most never will.

5.6 Addresses

Every engram has an address, `crystalline://engineering/retry-queue-gotcha:` the scheme, the domain, the permalink. It is the one absolute form. An agent hands it to another agent, `build_context` anchors on it, Fluid copies it to the clipboard in one click, and it survives a commit message, a ticket or a chat. A folder prefix appears in the permalink, `crystalline://engineering/queues/retry-queue-gotcha` if the engram had been filed under `queues/`, and a trailing `/` * anchors the whole folder at once.

The rule that trips people up is the scheme-less one. An identifier without `crystalline://` is domain-relative: `retry-queue-gotcha`, together with the domain named beside it, never `engineering/retry-queue-gotcha`, which would be read as a folder called `engineering` inside whatever domain you meant. The same holds for links: `[[Retry queue gotcha]]` resolves in the domain the linking engram lives in, and `[[engineering:Retry queue gotcha]]` reaches across. Either the address is absolute and carries its domain, or it is relative and the domain comes from context. There is no third form, and that strictness is what lets a permalink stay short.

5.7 The MANIFEST, up close

The last term is the first file in every domain. `MANIFEST.md` is an engram of type: `manifest` with the permalink `manifest`, which means it is read, indexed

and edited with the same tools as everything else. Its job is routing, and two sections do it. **## Scope** says what the domain covers. **## When to Use** says when an agent should route a task here, and those bullets, one line per domain, are what become the session-start briefing. **## Notes for Agents**, the third section the scaffold creates, holds the conventions the agent should follow inside this domain, the folder layout above all. Two optional sections come later. **## Tag Aliases** folds an old tag name into its canonical one so old searches keep working. **## Provisioning** names the folders of skills, slash commands, subagent definitions and MCP server configs a domain ships beside its knowledge. The tools that knowledge depends on travel with it, and install into a harness once a person says yes. That section gets a chapter of its own (Chapter 15). Tyrell's engineering manifest, a few months in:

```
---
type: manifest
title: engineering
permalink: manifest
tags:
  - manifest
status: stable
recorded_at: 2026-03-02
---

# engineering

## Scope

- The payments service, its retry queue and the services around it
- Deploy runbooks, team conventions and the decisions behind them

## When to Use

- Route questions about the payments service, the retry queue or a
  deploy here
- Check here before touching anything the team has already been burned
  by

## Notes for Agents

- Runbooks live under runbooks/, gotchas at the root
- One engram per gotcha; edit the existing one before writing a second
```


Write the When to Use bullets the way you would brief a new hire on which drawer holds what. They are read by a machine, but they are read the way a person reads a label.

Under the hood

The manifest being an engram is why routing needs no separate configuration: crystalline prompt system parses each registered domain's MANIFEST.md, takes the ## When to Use bullets, falls back to ## Scope when there are none, and emits one routing line per domain plus the behaviour rules. Two filenames in a domain are reserved and never engrams: index.md, the table of contents Crystalline re-generates after every write so the folder navigates itself on a git forge, and log.md, reserved but never generated. A title that would slug to either is refused.

In practice

In his second month Deckard starts the clients domain, because the team keeps re-explaining which customer is behind which integration. His first engram is about Weyland-Yutani's settlement batch, and it needs to say that the batch is exactly the thing that hits the retry cutoff. He types - depends_on [[Retry queue gotcha]], and the sweep tells him the link does not resolve, because clients has no engram by that title. He changes it to [[engineering:Retry queue gotcha]] and the edge appears, with a backlink on the engineering side that nobody on the engineering side had to write.

A week later Ripley reviews his pull request on the settlement job and wants the reasoning in the record. She pastes crystalline://clients/weylan-yutani-settlement-batch into the description. Her reviewer's agent reads that address, walks one hop into engineering and comes back with the cutoff, the bank rule and the runbook. Two domains, one edge and the string in the pull request is still readable in a year.

You now speak the language. The rest of the book is conversation.

In brief

A domain is a registered folder with a MANIFEST.md whose When to Use bullets route an agent to it; an engram is one markdown file with YAML frontmatter, prose, observations of the form - [category] text #tag and relations of the form - rel_type [[Target]], all on Google's Open Knowledge Format so nothing is locked in. status says whether knowledge holds, with stable as the default and deprecated, superseded, archived and legacy as the retirement set that fades in ranking without disappearing; salience is a rare, bounded lift. The absolute address is crystalline://domain/permalink, and anything without the scheme is

domain-relative. The manifest is itself an engram, which is why routing needs no configuration of its own.

READY. ■

PART III

Getting Started

CHAPTER 6

Installing locally

Don't Panic.

- Douglas Adams, *The Hitchhiker's Guide to the Galaxy*

Everything in this chapter fits in one sitting, and the sitting ends with a test: you will ask your agent about something it learned an hour earlier, in a session that has never heard of it, and it will answer. Until that happens nothing is installed, whatever the package manager says. So the chapter is built backwards from that moment. Install the binary, wire the harness, give the agent a domain: three steps, the same three on every platform, and then the test.

If you skipped Part II, two words carry the chapter: a domain is a registered folder with a `MANIFEST.md` that says when an agent should look there (Section 5.1), and an engram is one markdown file of knowledge inside it (Section 5.2). Everything below makes one of each.

6.1 Install the binary

Crystalline is one binary, and getting it onto a machine is whatever that machine's usual way is. On a Mac with Homebrew:

```
brew install jordiboehme/tap/crystalline
```

Debian, Ubuntu and their relatives get a `.deb` for amd64 and arm64 from the releases page, and it ships a systemd unit, installed disabled, that chapter 7 turns on. Windows gets an MSI that adds Crystalline to the `PATH` and upgrades in place; the Windows builds are not code signed yet, so expect a SmartScreen prompt and check the download against the `SHA256SUMS` file beside it. The bare binary ships for every platform too, and the macOS builds are signed and


notarized. The permanent address for all of it is on the colophon page: <https://github.com/jordiboehme/crystalline/releases/latest>.

Claude Desktop is the pleasant exception: no binary, no terminal. Download `crystalline-v<version>.mcpb` from the same page, one universal bundle for Apple Silicon Macs and Windows, with per-arch bundles for Intel Macs, then in Claude Desktop open Settings > Extensions > Advanced settings > Install Extension... and pick the file. Upgrading is the same motion, since Desktop has no in-app update for an extension: download the newer `.mcpb` and install it over the current one. Desktop users can skip the terminal parts of what follows; the agent creates its first domain itself.

6.2 First run

Nothing needs starting. The first agent that connects starts a background daemon, the one process that holds the index open, and every later agent, terminal and browser tab attaches to it. `crystalline serve` runs the same daemon in the foreground if you want to watch it, and `crystalline serve --daemon` in the background, but on a laptop most people never type either.

What you will type into is a browser. The daemon serves Fluid at `http://localhost:7411` by default, however it was started, and the first visit to an instance that has no accounts yet is a short form rather than a login page: pick a name and a password, and you are that instance's admin, already signed in (Figure 1). You are trusted to fill that form in because you are on the machine; a connection from the machine itself needs no token, and a bind that anyone else could reach gets a different treatment in chapter 7. Once any account exists the form is gone for good.



The screenshot shows a web browser window with the Crystalline Fluid interface. At the top is a green pixelated logo. Below it, the text reads "CRYSTALLINE Fluid" and "mind-meld your fluid thoughts with your AI's crystalline intelligence". A "Welcome" message states: "Nobody has an account on this instance yet. The one you create here is its first admin." Below this is a form with three input fields: "Name", "Password", and "Confirm password". At the bottom of the form is a green button labeled "Create admin account".

Figure 1: First visit: creating the admin account in the browser.

One more download belongs to first run. Semantic search wants the local embedding model; the daemon fetches it in the background on its first start, and crystalline model download fetches it ahead of time. Until then search is plain text, which works, and finds only the words you searched for.

6.3 Onboarding a harness

This is the step that turns a binary into a colleague, and it is one command:

```
crystalline install claude-code
```

Codex takes `crystalline install codex` and the Copilot CLI `crystalline install copilot`. What the command does deserves a paragraph, because every part of it corresponds to something a new hire is given on day one. It registers the MCP server with the harness, as `crystalline mcp --harness claude-code`, so the agent has the tools. It writes a session-start hook that runs `crystalline prompt system` on a fresh start, after a `/clear` and after a compaction, so every session opens with the briefing. It writes a stop hook, the once-per-session nudge from chapter 3 that asks the agent to capture what it learned before the session ends. And it copies the four skills into the harness's skills folder, the playbook that says when to use which tool. Run it twice and the second run

touches nothing. `--skip-mcp`, `--skip-hooks` and `--skip-skills` leave a part out, `--project` writes into the current repository's config instead of your global one, and `crystalline uninstall claude-code` reverses the lot while leaving anything that is not Crystalline's own alone.

The briefing is the part to look at once with your own eyes, because it is the whole of what your agent knows about you at the start of a session. `crystalline prompt system` prints it (Figure 2): an intro, the behaviour rules and one line per registered domain, drawn from the `## When to Use` bullets of that domain's manifest. A narrow question goes to the domain that owns it, a broad one sweeps all of them, and a write always names its domain. On a fresh install the domain lines are empty, which is correct: the desk is set up and nobody has shown the new hire the filing cabinets yet.

```

$ crystalline prompt system
CRYSTALLINE KNOWLEDGE ROUTING

You are being onboarded to the knowledge accumulated in this workspace. It is served by the
crystalline MCP server; these instructions govern its tools (your harness may prefix tool
names, for example mcp_crystalline_search_engrams). The domains below are registered and
ready to use; route searches and reads through them instead of starting from zero.

- engineering: Route questions about the payments service, the retry queue or a deploy here;
Check here before touching anything the team has already been burned by
- clients: Route anything about a customer, a contract or a renewal here; Check here before
preparing for a customer meeting
- board: Route anything about the funding round or a board meeting here

Behavior:
- Narrow question, one domain fits: search_engrams with domains=[that domain].
- Broad or unclear question: search_engrams with no domains sweeps them all.
- write_engram, edit_engram, move_engram and delete_engram always need an explicit domain.
- Pass write_engram a folder matching the domain's layout: a topic prefix build_context can
glob, a subfolder once a topic clusters, the root for singletons.
- Check vocabulary for tags and categories in use before inventing one.
- Salience (0-10) lifts an engram in search: set it at write time on exceptionally valuable
knowledge and raise it when one proves key to a task.
- build_context on a crystalline:// anchor gathers knowledge around a task.
- Read a MANIFEST with read_engram only when the domain's routing line is not enough;
list_domains with include_routing=true re-fetches the index mid-session.

```

Figure 2: The routing prompt an agent receives at session start.

Claude Desktop runs no hooks and needs none. Its server delivers the same routing block as part of the MCP server's instructions on every connection, and the optional companion skill from the releases page, uploaded under Settings > Capabilities > Skills, adds the capture and collaboration playbook on top. A chat client with no hooks that never shows the model a server's instructions gets a standing instruction instead: `crystalline prompt connector` prints the

paragraph to paste into its custom instructions, and the agent fetches its own onboarding with one `list_domains` call at the start of every session.

Under the hood

The install is recorded in a local receipt, and the receipt is what keeps a harness current: the first time a newer Crystalline runs, it refreshes the installed skills at session start, updating changed ones (an edited copy is kept beside the new one as `SKILL.md.bak`) and removing any the new version no longer ships. The `--harness` flag in the MCP registration is the other half of the same idea. A `crystalline mcp` process started with it asks the receipt whether that harness has its hooks wired, and if so serves that session neither the skills over MCP nor a second copy of the routing block, since the harness already has both as files. Nothing in that decision comes from the connecting client; it is settled before the session opens. A registration written by hand without the flag, or a remote client, is served the full surface, and `skills.serve` set explicitly to `true` or `false` makes every client identical.

6.4 The first domain and the first engram

Now the filing cabinet. It takes three prompts and no terminal: your agent creates the domain, labels it and writes the first thing into it, doing the scaffolding, the labelling and the writing through its own tools, `add_domain`, `edit_engram` on the manifest and `write_engram`:

```
You:    Set up a Crystalline domain called engineering in
        ~/knowledge/engineering. Its scope is the payments service:
        route questions about the payments service, the retry queue
        or a deploy there.
Agent:   Created and registered engineering. The manifest's Scope and
        When to Use sections carry what you just said, so your next
        session will see the domain in its briefing.

You:    Remember in engineering: the retry queue drops jobs older
        than 24 hours. It is a gotcha, tagged payments.
Agent:   Captured "Retry queue gotcha" in engineering, at
        ~/knowledge/engineering/retry-queue-gotcha.md.

You:    Why does the payments queue lose jobs?
Agent:   The retry queue drops jobs older than 24 hours, from
        "Retry queue gotcha" in engineering.
```


The embedding model is fetched automatically when the daemon starts, which the first agent to connect already did, so the last answer can come from the semantic ranking with no download step to remember. A Desktop user's path is exactly this, since there is no terminal there at all. It works even without the first prompt: an agent that finds nowhere to write creates a domain itself under Documents/Crystalline, and the first engram it captures lands there.

The bullets the agent wrote into the manifest deserve a second look. Open it before you go any further and make sure the `## Scope` and `## When to Use` sections say what you would say, because those bullets are the line your agent reads about this domain at the start of every session. Write them the way you would label a drawer for a new hire: "route questions about the payments service, the retry queue or a deploy here".

Under the hood, the first prompt was three commands. The quick start on the project's README does the whole thing in five lines, and the last three are these:

```
mkdir -p ~/knowledge/engineering
crystalline domain init ~/knowledge/engineering --name engineering
crystalline domain add engineering ~/knowledge/engineering
```

`domain init` scaffolds the folder with a starter `MANIFEST.md`; `domain add` registers it under the name the agent will use from now on, and indexes what is already there, which for the moment is the manifest. There is no separate sync step. The second prompt was one command, and running it by hand once shows the file appear:

```
crystalline write engineering "Retry queue gotcha" \
  --content "- [gotcha] The retry queue drops jobs older than 24h
#payments" \
  --tags gotcha,payments
```

Look in `~/knowledge/engineering` and there it is, `retry-queue-gotcha.md`, the file chapter 5 read line by line. Search it back with `crystalline search "retry queue"` and it is a plain text hit. On a machine where no daemon has fetched the model yet, fetch it, re-sync with embeddings and search for something the engram never says:

```
crystalline model download
crystalline sync --embed
crystalline search "why does the payments queue lose jobs"
```

The engram comes back from a question phrased nothing like it, which is the third prompt's answer, and the ranking your agent uses. Engrams written through Crystalline are indexed the moment they land; `crystalline sync` exists for files that arrive some other way, an editor or a git pull, when no daemon was watching. `crystalline status` shows what got indexed.

6.5 Verifying the loop

Now the test. Start a session in your harness and ask about the retry queue without explaining anything. Watch what the agent does before it answers: it searches engineering, because the briefing told it that is where the payments knowledge lives, and it answers from the engram, with the address in hand if you ask for it. That is recall, and it proves the wiring.

What that does not yet prove is the capture beat, which is the one that matters over months. Tell the agent something it could not have known, the way Ripley did on her Monday in March:

```
You:    The cutoff is 24 hours because Encom's chargeback window is
        three days - the queue needs two days to notice a job. Remember
        that, and why.
Agent:  Captured: added the chargeback window as the reason to
        "Retry queue gotcha" in engineering, with a depends_on relation
        to a new "Encom chargeback window" engram. Want me to adjust
        the wording?
```

Close the session. Open a new one, which has never heard of Encom, and ask why the cutoff is a day. If the answer names the chargeback window, you are done: the agent looked something up that it learned in a conversation that no longer exists. It is a Voight-Kampff test run in reverse. You are not checking whether the thing across the desk is a replicant; you are checking that it remembers.

When it fails, the fix is usually quick. If the agent answered from general knowledge without searching, the session did not receive the briefing: run `crystalline prompt system` yourself and check that engineering has a line, and if

it has none, the manifest's `## When to Use` section is empty. If the agent recalled the gotcha but not the reason, the capture never landed: look in `~/knowledge/engineering` for the file, and check that you said yes when it proposed writing. Once the test passes, close the terminal. Nothing in daily use needs it: the agent recalls and captures through its tools, the daemon starts itself and the one place you will go on purpose is the browser tab. The commands in this chapter were the installation. From here on, what you type is addressed to a colleague.

In practice

Deckard's first afternoon at Tyrell is the five lines above, on a laptop that arrived that morning. He runs the brew line and the install line, and instead of `domain init` Ripley has him check out the repository the team keeps its knowledge in and `domain add` that folder, so his agent's first session opens with a briefing that names `engineering` and everything the team has taught since March. He asks it about the retry queue to see what happens, and it tells him about the sin bin, the cutoff and the bank before he has read a single wiki page. He was onboarded in chapter 2 by a person over two days. His agent took a minute, and knew more.

Lambert, who runs the customer side and has never opened a terminal on purpose, installs the Desktop extension the same week. She has no domains, and her agent makes one the first time it needs somewhere to write, a `customers` folder under her Documents. By Friday it holds four engrams about the renewal she is working on, and she opens `localhost:7411` to read them. Two installs, one with a keyboard and one without, and the same loop running on both.

The desk is set up, the drawers are labelled and the new hire has remembered one thing overnight. Everything after this is a matter of what you choose to teach.

In brief

Installing is three steps on every platform: install the binary (Homebrew, a `.deb`, an MSI, or the Claude Desktop `.mcpb` that skips the terminal entirely and updates by installing the newer bundle over the old), wire the harness with `crystalline install <harness>`, and give the agent a domain with `domain init` and `domain add`. The install writes the MCP registration, the session-start hook that delivers the routing briefing, the stop hook that nudges capture and the four skills, and it is safe to run again. Nothing needs starting: the first agent to connect starts the daemon, which serves Fluid at `localhost:7411`, where the first visit creates the admin account. The install is finished when a fresh session recalls something a previous one learned.

Running a shared instance

The moment to leave localhost is the moment two people want to correct the same engram. Up to then a laptop is the right size: one person, one agent, one folder, nothing to operate. Past it, every fact the team learns is being taught twice and drifting once, and the fix is the oldest one in computing. Put it on a server, give everyone a login and let one instance be the one the whole team teaches.

The questions arrive in a fixed order, and the chapter follows it: where the instance runs; who may sign in, and how, the question chapter 4 handed here; how an agent on somebody's laptop proves who it acts for; and how one person keeps a domain to a smaller circle.

7.1 When to move off localhost

Chapter 1 settled where Crystalline runs: on a machine of yours, never on a hosted service. This chapter is about which machine. Three signals, any one of which is enough. Somebody who never opens a terminal wants to read what the agents learned and fix a fact without asking. An agent runs somewhere that is not a laptop, a hosted runner or a bot with a task, and needs to recall and capture like everyone else. Or the same engram has been corrected on two laptops in two directions.

A shared instance differs from a team domain, chapter 13's subject, where a folder lives in a GitHub repository, each person runs their own instance and knowledge travels between them as proposals the team reviews. A shared instance is one daemon everyone connects to: one index, one Fluid, one set of accounts, one folder of files on one disk. The two compose, and a mature team usually has both.

7.2 The shapes it takes

The full guide is docs/deployment.md in the repository, one scenario per section with a diagram each; this chapter summarizes. The shapes reduce to one picture (Figure 1): the daemon on a server, teammates' agents reaching it over MCP on HTTP instead of a local pipe, browsers reaching Fluid on the same port, and, where the team configures them, a proxy in front and an identity provider beside it.




Figure 1: One shared instance: agents connect over MCP with a token, people through Fluid with an account, with a proxy in front and an identity provider beside it where the team configures them.

The team server is a container. The published image binds `0.0.0.0:7411` inside it, keeps the knowledge as a bind mount so the files stay exactly the markdown they were, and puts the index on a volume:

```
docker run -d \
  --name crystalline \
  -p 7411:7411 \
  -v "$(pwd)/knowledge:/knowledge" \
  -v crystalline-data:/data \
  ghcr.io/jordiboehme/crystalline:latest
```

The `with-model` tag bakes the embedding model in, for a host with no egress. The same shape without a container is the `.deb`'s `systemd` unit: put `CRYSTALLINE_SERVICE_HTTP=0.0.0.0:7411` in `/etc/default/crystalline` and run `sudo systemctl enable --now crystalline`. A scale-out variant serves Fluid from an `nginx` tier in front of the daemon, and a read-only shape, `--read-only`, serves knowledge curated as a reviewed repository: the writing tools vanish from what an agent is offered while the index keeps following every pull.

Whatever shape you pick, a few things follow you into it. An agent reaching the daemon by a hostname needs that name in `CRYSTALLINE_SERVICE_ALLOWED_HOSTS`, since the MCP transport checks the `Host` header against a loopback-only allow-list to block DNS rebinding. Serve it over TLS anywhere but `localhost`, since the session cookie is marked `Secure` off loopback and a browser on plain `http://team.lan` will not keep it. And a proxy in front must forward `Host` verbatim, pass the `WebSocket` upgrade through for co-editing and accept 10 MiB request bodies, as the daemon does.

7.3 Accounts and roles

Nothing is readable until an account exists. A fresh instance greets the first browser with the `create-first-admin` form from chapter 6, with one difference: a daemon bound to anything but loopback cannot trust a visitor by where they connected from, so it prints a one-time setup token as it starts and the form asks for it. Read it in `docker logs crystalline`, or `journalctl -u crystalline` under `systemd`; a restarted daemon prints a fresh one. The terminal route needs no token: `crystalline users add ada --role admin` writes the accounts database directly, and a running daemon picks the account up on its next lookup.

Three roles, each containing the one before: a viewer searches, reads and browses; an editor also writes and edits engrams; an admin also manages domains and users, from `crystalline users` or from Fluid. Accounts live in a small database beside the index, never inside it, so `reindex --full` cannot take them along; on the container it sits on the `/data` volume and no `rebuild` restores it, so back the volume up.

When somebody leaves, an admin disables or deletes the account, and deleting sweeps whatever the account had connected. That is the whole deprovisioning story; there is no directory sync and no SAML.

7.4 Single sign-on

Local accounts are right for a team of nine. Inside a company that already has an identity provider, people should sign in with the login they already have, and Crystalline talks to the provider through one generic integration, OpenID Connect, with no adapter per vendor. The configuration is three keys, `auth.oidc.issuer`, `auth.oidc.client_id` and `auth.oidc.client_secret`, plus a label for the button and the scopes if the defaults do not suit; discovery, PKCE and key rotation are the integration's job. Fluid's login page then shows "Sign in with" the label beside the local form (Figure 2), and both keep working, so an instance can run mixed: local accounts for bots and the break-glass admin, everybody else through the provider.




Figure 2: Signing in through the instance's identity provider.

The first successful sign-in creates the account, named from the provider's preferred username and made unique if it collides, at the role `auth.oidc.default_role` says, viewer unless you change it. What the account is keyed on is what most such integrations get wrong: the durable key is the pair of issuer and subject, the provider's own stable identifier, never the username, the email or the display name. Those are presentation; rename yourself at the provider and you land in the same account with a new label. And there is no silent linking: signing in through the provider with an email that matches an existing local account gets you a fresh account, never the existing one, and

linking the two is an explicit act by the signed-in local user or by an admin. The provider authenticates and Crystalline authorizes; a group at the provider does not become a role here.

Entra ID: use the tenant-specific issuer

Microsoft Entra publishes a tenant-independent discovery document whose issuer field carries the literal text {tenantid}, and an issuer check against that template fails for every token. Crystalline refuses that value at configuration time and names the fix: the issuer must be your tenant's own `https://login.microsoftonline.com/<your-tenant-id>/v2.0`. Prefer a certificate credential over a client secret where the deployment allows it. And know that Entra's subject identifier is pairwise per app registration, so re-registering the application breaks every link between provider identities and Crystalline accounts, and an admin has to re-link them.

7.5 Trusted-header mode

Some deployments already have a proxy that authenticates everybody and names the arrival in a request header, and Crystalline can consume that. `auth.trusted_header` names a single custom header, `remote-user` say, and an account it names is created at viewer role the first time it is seen.

A second form, `auth.proxy_headers`, takes the standard set that forward-auth proxies emit, `Remote-User`, `Remote-Name`, `Remote-Email` and `Remote-Groups`, using the first for identity, the next two as presentation and reading the groups without acting on them. Pick one form or the other; a daemon configured with both refuses to start. Both are off by default, and a header that arrives while the mode is off buys the sender nothing.

Now the warning, which the documentation does not soften, so neither does this page. Header mode trusts the proxy completely, because a header is a line of text anyone can add to a request. It is safe only when Crystalline is unreachable except through the proxy, and the proxy strips any client-supplied copy of those headers before setting its own. A daemon reachable on its port directly is an instance where anybody can be anybody. If you cannot say with confidence that the only path to the port runs through the proxy, use single sign-on instead. And create the first admin before the proxy goes in front: a request carrying forwarding headers is never counted as local, so the setup form behind a proxy asks for a token that a loopback-bound daemon never printed.

7.6 Agent authentication

On a laptop the MCP endpoint is open, on loopback, to whatever local process reaches it, and that is the right trade for one person. On a shared instance turn on `auth.mcp`, and from then on every agent that connects over HTTP identifies itself, exactly as a person does. There is no half-way tier: a connection without credentials is refused before any session opens, with a message naming how to get one; HAL refused Dave with less courtesy. The open tier exists only on an instance that leaves the setting off, and sees only the domains everyone sees. An agent that could search a shared instance anonymously would be a colleague with no name on the commit.

The credential is a token that belongs to a person. Each user issues their own in Fluid, under the Agent access card on their profile, or an admin mints one for any account with crystalline users `mcp-token <name>`; either way it is shown once, and the instance keeps only a hash. It travels as an `Authorization: Bearer` header on the server's entry in the harness's MCP registration, never expires on its own and stops working the instant it is revoked or rotated, because every request looks it up. Every account may issue one, viewers included, and that is safe because of what the token means.

What it means is that the agent is its user. It sees what that person sees, writes where that person may write, shares under that person's identity when a team domain is involved, and the engrams it captures record which client acted for which account. A viewer's agent is therefore read-only with nothing further to configure, and a leaked token names one person, whose one token is revoked. One caller needs no token: the CLI and a stdio session on the server itself are the machine's owner and see everything, which is the honest statement of what a filesystem already permits.

7.7 Private domains

The first rule of a private domain is: you do not talk about the private domain. The second rule of a private domain is: the instance DOES NOT talk about the private domain.

- with sincere apologies to Chuck Palahniuk

By default whoever may sign in sees every domain. Some knowledge belongs to a smaller circle, the founders' notes on a funding round or one person's working knowledge on a company-wide instance, and a domain can be made private.

It is then visible to its owner, to the instance's admins and to whoever the owner invites, each member at one of three levels: a viewer reads and searches, an editor also writes and a manager also runs the membership. The owner is one person, and only the owner or an admin changes the domain's visibility, transfers it or deletes it. Transfer is how a private domain survives its owner leaving.

Two truths belong here, and the second is printed on the toggle itself so nobody discovers it later. Instance admins see and administer every domain, private ones included. And a private domain protects its contents from the instance's other users, not from whoever operates the machine, because the engrams are files on that machine's disk, with no encryption at rest. If the threat is the operator, the answer is a different machine.

Visibility follows the caller onto every surface: search, reading, `build_context`, the routing briefing, the graph, the maintenance sweep, exports and Fluid's domain list all show a private domain only to those who may see it. To everyone else it does not exist, literally: a search or a read against a domain you may not see answers exactly as it would for one that was never registered, because a "forbidden" would already say something. That is the second rule. Writes check the membership level rather than the instance role, so an instance editor invited as a viewer is told, when they try, that editor access is what the domain requires. Members are managed from a Members card in the domain's settings, or on the server with `crystalline domain members <domain> add <user> --level editor` and its `list` and `remove` siblings. Agents inherit all of this from their user and administer none of it; there is no tool for changing membership, on purpose. Membership records live in the accounts database beside the index, so a wipe and resync of the index leaves every permission where it was.

In practice

Tyrell moves in June, after the second time Ripley and Deckard corrected the same engram in two directions. Ripley runs the container on the team's server, behind the TLS proxy that fronts everything else, the first admin created before the proxy went in. She points `auth.oidc.issuer` at the company's identity provider, everybody signs in on Monday with the login they already have, and she promotes the engineers to editor. Then she turns on `auth.mcp`. Deckard's agent, pointed at the server over HTTP, is refused at once with a message saying where to get a token; he issues one on his profile, pastes it into his harness registration, and the next session opens with a briefing that names engineering and clients, learned by everyone, recalled by his.

Dallas, who runs the company, keeps a board domain for the funding round. He makes it private and invites Ripley as a manager and Lambert as a viewer. Deckard's agent, asked about the round, searches every domain it can see and reports that nothing is known; Ripley's, acting as Ripley, finds the notes at once; Lambert's, trying to capture a fact into board, is told that her membership is viewer and proposes writing to clients instead. One team, one instance, one folder of files on one disk, and every action in it with a name attached.

Somebody has to be the one who runs the server, and the rest of the team gets to forget it exists. That is the whole point of a shared instance: one place everybody teaches, and one login that says who did.

In brief

Move off localhost when a team wants to teach one knowledge base: run the container, the systemd unit or the read-only shape from docs/deployment.md, serve it over TLS, and create the first admin with the setup token the daemon prints or with `crystalline users add`. Accounts come in three roles, viewer, editor and admin, created locally or on first sign-in through any OpenID Connect provider, keyed on issuer and subject and never linked by email; trusted-header mode is safe only when the proxy is the only path in. With `auth.mcp` on, every agent over HTTP authenticates with a per-user token and acts as that user, with no unauthenticated tier beside it. A private domain is visible to its owner, the instance's admins and invited members on every surface, and it protects from other users, never from whoever operates the machine.

READY. ■

PART IV

Working with Crys- talline

Onboarding an agent

On a new colleague's first morning somebody hands them the floor plan. Payments is on the second floor, the people who know about the bank sit by the window, the runbooks are in the cupboard that sticks. Which door leads where is the whole of day one, and it is enough, because a person who knows where things are can find out everything else themselves.

That floor plan is what an agent receives at the start of every session with Crystalline. This chapter is about the day it arrives: what the briefing says, how the agent finds its way with it, what it means to hand a colleague a task instead of a question, and the reminder at the end of the day that turns a session into experience. Chapter 6 showed the wiring. This is what the wiring is for.

8.1 The briefing at session start

A session opens, and before the agent has read a word of your request, the session-start hook runs `crystalline prompt system` and puts its output in front of the model. Here is Tyrell's, a few months in, abridged:

CRYSTALLINE KNOWLEDGE ROUTING

You are being onboarded to the knowledge accumulated in this workspace. It is served by the crystalline MCP server; these instructions govern its tools. [...]

- engineering: Route questions about the payments service, the retry queue or a deploy here; Check here before touching anything the team has already been burned by
- clients: Route anything about a customer, a contract or a renewal here; Check here before preparing for a customer meeting

Behavior:

- Narrow question, one domain fits: `search_engrams` with `domains=[that`

```

domain].
- Broad or unclear question: search_engrams with no domains sweeps
  them all.
- Read a MANIFEST with read_engram only when the domain's routing line
  is not enough; list_domains with include_routing=true re-fetches the
  index mid-session.

```

Read it as the agent reads it. The opening says what this is and whose tools it governs, so an agent with four MCP servers connected knows which `search_engrams` is meant. Then one line per domain, and every word of those lines came from a person: they are the `##` When to Use bullets of each manifest, joined with semicolons, exactly as the team wrote them (Section 5.1). Then the behaviour rules, the same on every instance, naming the exact tool each one reaches for. That is the floor plan (Figure 1). It says where things are, nothing about what is in them, and the agent goes and looks.

```

$ crystalline prompt system
CRYSTALLINE KNOWLEDGE ROUTING

You are being onboarded to the knowledge accumulated in this workspace. It is served by the
crystalline MCP server; these instructions govern its tools (your harness may prefix tool
names, for example mcp_crystalline_search_engrams). The domains below are registered and
ready to use; route searches and reads through them instead of starting from zero.

- engineering: Route questions about the payments service, the retry queue or a deploy here;
Check here before touching anything the team has already been burned by
- clients: Route anything about a customer, a contract or a renewal here; Check here before
preparing for a customer meeting
- board: Route anything about the funding round or a board meeting here

Behavior:
- Narrow question, one domain fits: search_engrams with domains=[that domain].
- Broad or unclear question: search_engrams with no domains sweeps them all.
- write_engram, edit_engram, move_engram and delete_engram always need an explicit domain.
- Pass write_engram a folder matching the domain's layout: a topic prefix build_context can
glob, a subfolder once a topic clusters, the root for singletons.
- Check vocabulary for tags and categories in use before inventing one.
- Salience (0-10) lifts an engram in search: set it at write time on exceptionally valuable
knowledge and raise it when one proves key to a task.
- build_context on a crystalline:// anchor gathers knowledge around a task.
- Read a MANIFEST with read_engram only when the domain's routing line is not enough;
list_domains with include_routing=true re-fetches the index mid-session.

```

Figure 1: The same block unabridged, as `crystalline prompt system` prints it.

How it reaches the desk depends on the harness, as chapter 6 said. Claude Code and Codex run the hook on a fresh start, after a `/clear` and after a compaction, and skip a resumed session; the Copilot CLI's hook does the same on a fresh

start and stays quiet on a resume; Claude Desktop gets the block inside the MCP server's own instructions; a remote client fetches it with one `list_domains` call. Either way the agent starts out knowing which drawer holds what and nothing more, the state a good new hire is in at nine on Monday.

Under the hood

The block is built to survive being cut off. A harness in tool-search mode truncates a server's instructions at roughly two kilobytes, so the MCP rendering keeps under a 1900-byte budget and degrades in fixed steps: every routing line drops to one bullet, then the list collapses to a count line pointing at the re-fetch call.

8.2 How the agent finds its way

A routing line is a targeting aid. The routing skill says so in as many words, and adds what every colleague knows about floor plans: the label on the door is not the inventory of the room.

The first decision is scope. A narrow question, one that clearly belongs to one service, team or customer, goes to the domain whose line describes it, in a single `search_engrams` call with `domains` set. A broad or unclear question omits `domains` and sweeps everything, each hit labelled with its domain. The agent does not stop to ask you which it is: a search is cheap and usually settles the ambiguity, and a clarifying question is for after it came back empty.

The second decision is when to widen. A scoped search can return plausible near-misses when the subject belongs to one domain and the question to another: who gets paged about refunds is a people question, not a payments one. When the obvious domain keeps missing after a rephrasing or two, the agent runs one unscoped sweep. Then it reads. A snippet is a short window around a match, and a colleague who answers from a snippet is guessing, so anything that summarizes or quotes an engram goes through `read_gram` first. A strong hit becomes an anchor for `build_context`, which pulls in what the engram depends on before the work starts.

The agent was never given the engrams, only the doors, and what is behind each is fetched when a task calls for it. A colleague with a floor plan beats one carrying the building around in a wheelbarrow.

8.3 Assign, do not copy

You would not message a colleague a question, copy their answer into your editor and do the work yourself. You would assign them the task, because doing it is how they get better at it. An agent with Crystalline’s tools is in the same position, and treating it as a lookup service wastes most of what you set up.

Give it the task. “Add a metric to the retry queue” lands better than “what do you know about the retry queue”, because the first sentence sends the agent through the whole loop on its own: it routes to engineering, recalls the gotcha and the bank rule, reads the engram in full, walks one hop to the deploy runbook, writes the change in Python and tells you what it found on the way. The floor plan only pays off when somebody walks the corridors, and the somebody should be the agent.

It follows the processes a person would, since it can be wrong in the same ways: review for its changes, a yes for its captures. A colleague you never delegate to never learns your work, and neither does an agent.

In practice

Deckard’s first week with his agent, session by session. Monday is the afternoon from chapter 6: the briefing names engineering, and the agent tells him about the sin bin, the cutoff and Encom. On Tuesday he is given a real task, changing the alert threshold on the queue, and explains nothing. The agent routes to engineering, recalls the gotcha, finds the config file Ripley’s session captured earlier that week, and makes the change. On Wednesday it is wrong for the first time: it proposes rolling out on Friday afternoon, and Deckard tells it the team never deploys payments on a Friday. It asks whether to capture that in engineering as a [convention]. He says yes.

Thursday is a long session on the deploy runbook, and when it ends the reminder fires once: the agent proposes two captures, a gotcha about a stale health check and a lesson about rollout order, and raises the salience of the retry-queue engram. On Friday he asks it to schedule the rollout, and it answers, before he has mentioned a day, that Friday is out by the team’s convention and suggests Monday. On the first session it knew where the drawers were. On the fifth it told him something about his own team, unasked, that he had taught it two days before. That is the moment a stranger stops being one.

8.4 The reminder at the end of the day

A first week only counts if what was learned survives the weekend, and that is the beat agents skip. Left alone, a model finishes the task and stops, and

the gotcha it met at minute forty goes with it. So crystalline install wires a reminder, the closest thing the loop has to a habit.

It is a Stop hook. On the first stop after a session has gained real substance it asks the agent to review the conversation for durable learnings: facts, decisions, patterns, gotchas, corrections from you. For each one not yet captured the agent proposes an engram in the fitting domain and waits for a yes; a correction that makes an existing engram wrong becomes a proposed edit or supersession rather than a second engram. If a recalled engram proved to be the key to the task, the agent raises its salience. If nothing qualifies, it finishes without mentioning the check. The reminder fires at most once per session, and costs about a hundred and twenty tokens. When maintenance is due, a paragraph rides along asking for a sweep; when a team domain holds unshared work, another asks the agent to propose sharing, still a proposal, because sharing puts somebody's work in front of reviewers.

An agent that reaches a shared instance only over MCP has no hooks. Crystalline's one reliable channel to a model mid-session is a tool result, so the reminder rides there: a write, edit, move or delete comes back with a single clearly prefixed line appended when something is due, never a search or a read, so recalled knowledge is never mixed with guidance. Sharing and maintenance travel this way. Capture does not, because it is a reflection at the end of a conversation and the server cannot tell when a conversation ends; the routing block keeps teaching capture as standing practice instead. Each person is nudged on their own cadence, a few hours apart:

```
[crystalline] 3 changes across engineering are not yet shared. If  
that work is done, propose sharing with share_changes so the domain  
owner can review it - and wait for a yes.
```

The floor plan gets a new hire through the door. What they carry back at the end of the day is what makes them worth having on Tuesday.

In brief

Every session opens with a routing block: one line per domain from the manifests' When to Use bullets, and the behaviour rules naming the exact tools. Each harness delivers it the way chapter 6 wired it. The agent scopes a narrow question to one domain, sweeps for a broad one, widens when scoped hits keep miss-

ing and reads before it answers. Hand it tasks rather than questions, and let the once-per-session reminder close the day: propose captures, wait for a yes.

CHAPTER 9

Teaching

The best onboarding you ever had was a walk. Somebody took you across the floor, stopped at a machine and said “that one lies”, and you never trusted it again. Nobody wrote that down. It was not on the slide. It was the kind of knowledge that only moves from one head to another when a person decides it matters enough to point at, and the pointing is what this chapter is about.

Day one, last chapter, was the floor plan. Now the proper onboarding: what you tell a colleague in the first month, how you tell an agent the same things, what a well-taught fact looks like on disk, and the two rooms where the teaching happens, a session and a browser tab. It ends with two people editing the same page at once, which is where teaching stops being a lecture.

9.1 What to teach

Ask any good manager what a new hire needs and the list comes out in the same shape every time. What we value, and what we will not do even when it would be faster. How things are done here: the review, the deploy, the meeting that actually decides. What is expected of the role. Which tools we use, and which we abandoned and why. And where information lives: which repository, which dashboard, which person. An agent needs the same list, in the same order of importance, and the last item is the one that makes the other four cheap.

A person’s knowledge does not look like a source tree. No engineer on your team holds the payments service in their head by the letter. They know it is in the payments repository, that the retry logic is under queue/, that the config nobody documented is in the deploy chart, that Ripley owns it and Encom’s rule explains its strangest number. They know how to get access and which door to knock on. What they carry is a map of the territory plus the concepts to navigate it fast, and that is exactly what to teach an agent: not the source tree, which it can read with the tools its harness gives it whenever a task calls for it, but which repository holds the code for which purpose, how access is granted,

the architecture and the vocabulary. Pointer knowledge, access knowledge and the mental model. Teach the map and the territory stays where it is, reachable.

So the engrams that earn their place in the first month are the ones a person would say across a desk. The chargeback rule behind the cutoff. The name the team uses for the queue. The fact that the wiki says gateway where the code says edge. The rule about Fridays. Not one of them is on any page, and every one of them is what an experienced colleague knows.

9.2 Teaching through the agent

The simplest way to teach is to say something worth keeping, in a session, and watch what happens.

Tell the agent “the settlement job runs at six because Encom’s file lands at five, and it fails silently if the file is late”, and it does not write. It searches first, in engineering, because the capture skill makes search-before-write a hard gate, and it checks whether an engram already owns the topic, in which case the new fact becomes an edit of that engram rather than a second one beside it. Finding nothing, it proposes: “Should I capture this in engineering as a [gotcha]?” A terse “store this” is already a complete instruction and skips the question; an aside you obviously did not mean to keep gets no question at all. It names the domain every time, since a write has no default, and when the owner is not obvious it asks rather than guesses. A rotation detail learned during a payments incident is people knowledge, and the agent sweeps every domain before filing it where the incident happened to be.

Then it writes a well-formed file: a title, tags drawn from the vocabulary already in use rather than a near-synonym, a folder matching the domain’s layout, and a body of at least three lines, because verification rejects a lone bullet as too thin. Hand it one sentence and it pads honestly, one bullet for the fact in your words, one for how and when it was learned, one for the scope and never an invented specific. You glance at the result the way you glance at a new hire’s first commit, and say yes.

Correction is the same conversation. “The file lands at half past five, not five” finds the engram and changes the value where it stands. “We moved settlement to the new broker in August” is a different thing, the world changing, and the agent supersedes rather than edits, as chapter 3 walked through. Teaching and correcting are one motion, which is why the framing in chapter 2 said teach, rather than file.

9.3 Teaching through Fluid

The other room is the browser. Sometimes the fastest way to teach is to write it yourself, without a model in the loop and without spending a token: open the engram in Fluid, switch to edit and type (Figure 1). The editor previews live, the frontmatter is a form, and typing `[[` completes titles across every domain, so a relation is two keystrokes and a pick. The file on disk is what changes; the page you were reading (Figure 4) is that file.




Figure 1: Editing an engram in Fluid: live preview, the frontmatter form and wikilink completion.

What you write there does not sit unreviewed. The engram records that a person wrote it, `human:` and your account name in its `generated.by`, and the next maintenance sweep raises it as a finding of its own: captured by a person and never reviewed. The finding is judgment class, so the agent does not rewrite your words. It reads the engram, verifies the claim, aligns the tags with the vocabulary, wires relations both ways into the neighbourhood the engram belongs to and records a verified entry; any change to the wording it proposes and waits. Chapter 2's Deckard corrected a stale page by hand, and that Friday's sweep is where the agent tidied it into the graph. That is the division of

labour: you teach in whichever room is nearest, and the agent keeps the archive coherent.

9.4 The anatomy of a good engram

Chapter 5 read an engram line by line. What makes one good deserves a page of its own, because the difference between a good engram and a page is the difference between a colleague who explains and one who forwards you a link.

One unit of knowledge. An engram is about one thing, and when the thing grows the engram is edited, not forked. Prose first: the explanation you would give across a desk, what the thing is and why it is the way it is, because prose is what search matches and what a reader reads. Then observations for the facts, one bullet each with the most precise category you have, [fact] for verified state, [decision] for a choice, [gotcha] for the trap, [lesson] for what cost an hour, [risk] for the concern, never [fact] for a guess. Then relations for structure, depends_on, supersedes or a plain relates_to, so the graph can carry a reader to what the fact leans on. Tags from the vocabulary in use. A salience only when the engram is genuinely exceptional, the rule a dozen other rules hang from, since most never need one. And an engram is what is true now: a corrected value is replaced where it stands, never stapled on as a dated update, and the validity dates stay absent unless a fact is genuinely bounded. This one Deckard's agent wrote into clients, after the vignette below:

```
---
type: engram
title: "Weyland-Yutani CT0: one-pagers, decides live"
permalink: weyland-yutani-cto-one-pagers-decides-live
tags:
- weyland-yutani
- meetings
status: stable
recorded_at: 2026-06-11
generated: { by: claude-code/2.1.233, at: 2026-06-11T17:05:41+00:00 }
---
```

Vickers, Weyland-Yutani's CT0, skims decks and decides in the room. Of this quarter's three meetings, the two that landed led with a one-pager and the numbers; the one that went sideways led with a deck.

```
- [pattern] Lead with a one-pager and the numbers #weyland-yutani
#meetings
```

```
- [gotcha] A deck loses her by slide three; decisions happen live  
#weyland-yutani  
- [fact] Derived from the notes of three meetings, April to June  
#meetings  
- relates_to [[Weyland-Yutani settlement batch]]  
- summarizes [[Weyland-Yutani meeting notes, Q2]]
```

Notice what is in it and what is not. The pattern, the trap and the provenance are there, and the notes are kept beside it, attached in the agent's word and on disk as a type: source engram under sources/ linked both ways, so the next reader can check the derivation. The minutes of the three meetings are not in the engram itself; a summary is what most recalls read, and the full text is pulled in only when the summary falls short.

In practice

The Weyland-Yutani integration is Deckard's, which puts him in the room for the quarter's meetings, and after the third one he sits down with his agent:

After the last meeting

You: That went sideways. Here are my notes from the last three Weyland-Yutani

meetings - learn how their CT0 actually decides.

Agent: Learned them. The pattern: she skims decks, decides live in the

room, one-pagers land. Captured "Weyland-Yutani CT0: one-pagers, decides

live" into clients (#weyland-yutani #meetings), notes attached.

Next quarter, a fresh session

You: Prep me for Thursday's Weyland-Yutani renewal.

Agent: One-pager coming up - recalled that their CT0 decides in the

room and decks land badly. Leading with the renewal numbers.

Look at what he did not do. He did not tell the agent a fact. He handed it the source material and a question, the way you hand a colleague the notes and say "work out how she decides", and the agent derived the pattern, wrote it down with the notes kept beside it as provenance, and three months later recalled it in a session that had never heard of Vickers. That is teaching at its most economi-

cal: you supply the material and the judgment about what matters, and the colleague does the learning.

9.5 Attachments

Some teaching material is a file. The chargeback agreement, the deck a vendor sent, the diagram you drew at the whiteboard and photographed. In Fluid you paste, drag or upload it onto the engram it belongs to, and it lands under the domain's assets/ folder with a reference in the engram's body, an image `![...](assets/diagram.png)` or a link `[...](assets/deck.pdf)`. Images, PDFs, slide decks, office documents and data files are all accepted, up to ten megabytes each. Files enter through Fluid or a domain archive, never through an agent write, and an attachment always belongs to some engram's teaching: a file nothing references comes up in the sweep as something to delete or to turn into knowledge.

They are not opaque to the agent. `read_engram` returns each attachment as a resource link, and the agent fetches the file itself over MCP before summarizing what it shows. The sweep closes the loop from the other side: an attachment nobody has analyzed is a finding, and working it means reading the file, capturing what it teaches and marking on the engram which file was analyzed at which version, so a changed deck comes back around. One small courtesy for images: a formatting fragment on the target, `![Chart](assets/chart.png#right,w=50%)`, floats and sizes the picture in Fluid, and every other markdown renderer ignores it.

9.6 Editing together

Chapter 3 said the way a colleague really learns is by pairing, and Fluid makes pairing literal. Open an engram that somebody else has open and you see their cursor; type, and both sets of edits merge conflict-free and land as one save. Two people, one page, the whiteboard version of a wiki.

On a shared instance where agents authenticate (Chapter 7), the third name in the strip can be the agent's. Its user has the engram open; the agent, acting as that user, edits the same engram from the session, and the two are attached to one live document. The agent appears in the participant strip as a named peer, the account with an agent designation and a colour of its own, and its edits highlight in that colour. An edit it makes is computed against the page

as it stands at that moment, so it composes with what you typed seconds ago instead of stomping it; a read it makes sees your unsaved typing rather than a stale copy; and its write receipt says the change landed in a live page and names who is present, so it can tell you “I updated the section you are looking at” and mean it. Only a wholesale overwrite of a page you have open asks first. That is the fluid and crystalline halves working on one page at once, human with human and human with agent, and chapter 13 puts a whole meeting in that room.

Teaching a colleague is pointing at the machine that lies. Teaching an agent is the same gesture, and this one writes it down.

In brief

Teach an agent what you would teach a new hire: values, processes, expectations, tools and above all where information lives, taught as a map rather than the territory, which repository holds what, how access works, the architecture. In a session, say something worth keeping and the agent searches first, proposes a capture in the fitting domain and writes a well-formed engram once you say yes, with corrections edited in place and changed facts superseded. In Fluid, write it yourself, and the next sweep verifies, tags and wires what a person captured without rewriting their words. A good engram is one unit of knowledge with prose, categorized observations, typed relations and, rarely, a salience; attachments hand over files that agents read back. Co-editing merges live, and an authenticated agent can sit in that session as a named peer.

Ingesting knowledge at scale

Nobody has ever onboarded a colleague by pointing at four hundred wiki pages and walking away. Plenty of people have, and the colleague spent a month believing the architecture diagram from the year before the rewrite and calling the product by the name it had until the acquisition. The pages were not wrong when they were written. They were simply handed over without the ten minutes at the whiteboard that tell a newcomer which of them to believe.

Those ten minutes are the subject of this chapter. Sooner or later the teaching of chapter 9, one fact at a time, meets the existing pile: the wiki, the shared drive, the decks. The agent can learn from it in an afternoon, and it wants to, like Johnny 5 flipping through pages of an encyclopedia: input, more input, need input. What it cannot do is bring the judgment a person brings to a wiki, the sense of what is current, what was aspirational and which of two contradictory pages to trust. Chapter 1 said that judgment has to be handed over. This is how, and how you check what came back.

10.1 What a briefing contains

A briefing is the paragraph you would say to a person before they opened the wiki, and the useful ones are concrete about the past. Names first: the product now called Settlement was called Ledger until the reorganization two years ago, so normalize to Settlement; the wiki says gateway where the code says edge, and edge is the word to keep. Then a rule for contradiction, since a wiki of any age contains them: trust the newer page, unless one of them is the runbook. Then what to skip: the archived spaces, the personal pages, anything under Drafts. Then the one everybody forgets because it is obvious: the decks were written to persuade and the wiki is what we actually do, so where they disagree the page wins and the deck's version goes in as an [idea] at most.

Add where things go: which domain, which folder layout, the tags in use. And the shape of the review you want: propose before writing, keep source links, ask about anything resembling a decision nobody wrote down. At Tyrell it fits in one message:

```
You: Read the engineering space of the wiki through the connector and
      propose engrams for the engineering domain. Ledger is now called
      Settlement - normalize. Where the wiki says gateway the code says
      edge; use edge. Skip the Archive space and anything under Drafts.
      On a contradiction trust the newer page unless one is a runbook.
      The Q1 platform deck is aspirational; the wiki is what we do.
      Keep a link back to every source page and its last-edited date.
      Propose the list first and I will trim it.
```

That message is the difference between an ingest that lands and one that produces a confident, well-formatted archive of things that stopped being true.

10.2 The sources, and who reads them

Crystalline ships no scraper, by design. The agent is the ingester, and it reads sources with whatever its harness gives it: the wiki over the MCP connector the harness provides, a folder of exported documents through its file tools, a repository through the code tools it already has. A repository, though, is a source the agent consults rather than one it ingests: the repository holds the code, and the engrams hold which repository holds the code for which purpose, how to get access and the architecture and concepts that make the right file findable in a minute. The division is precise: the harness supplies the reading, the agent supplies the distillation and Crystalline is the destination, where what was derived lands as engrams with provenance, and the review loop that checks it afterwards.

A wiki is the common source, read page by page through the connector. A folder of documents, PDFs and exports and the odd spreadsheet, is read the same way. Slide decks are the one source that usually comes in through Fluid instead: a person drops the deck onto the engram it belongs to, and the agent reads it back over MCP as an attachment, so the deck becomes teaching material rather than a file nobody opened (chapter 9). And a tree of markdown that already carries frontmatter, an old knowledge base that was nearly an archive already, does not go through the agent at all: `crystalline import <src> --domain <name>` folds

it into a domain as a pure file transformation, dry run first, normalizing legacy type values, backfilling status and dropping sentinel dates on the way.

10.3 Distill, then propose

The cardinal sin of ingestion is mirroring: copying a source wholesale so that the archive looks complete, then filling the gaps with guesswork. The playbook that ships with Crystalline uses the right example. A vendor recall notice is forty screens of marketing header, revision table, legalese and support matrix, with three paragraphs that matter. Distilled, it becomes two observations, a resource link back to the page and a `source_date` in the frontmatter. Forty screens, two facts, one link.

Distilling is also where the briefing does its work. The agent reads a page about Ledger and writes Settlement; it finds a deck slide claiming a feature and files it as an [idea] because you told it decks persuade; it meets two pages that disagree and applies the rule rather than picking one silently. And one shape earns verbatim retention, only ever beside its distillation: a document whose exact words you may need to quote. The summary stays an ordinary engram where its topic lives; the full text goes in as a `type: source` engram under `sources/`, linked both ways with `summarizes` and `summarized_by`, split into parts when it outgrows the verify budget. On recall the summary is what gets read, and the full text is pulled in only when the summary falls short.

Then, before anything lands, the list:

```
Proposed for engineering (23 engrams):
runbooks/ Settlement deploy runbook
runbooks/ Retry queue drain procedure
          Edge routing rules
          Encom file format, v3
...
Write these? (y / n)
```

You trim it. Proposing first is the rule for a whole space, not just a page, because the trimming is where your judgment enters the archive, and striking five lines off a list is cheaper than retiring five engrams next month.

10.4 The review loop

What the agent writes is a proposal in a second sense: it carries provenance, and a person reviews it. Each engram names its source, a link, a date, an attachment, so the reviewer can open the page it came from and check the derivation. The review happens in Fluid, where an engram is a page and a fix is an edit, and a reviewer who knows the material corrects the one value the agent got wrong without going back through a session. A team domain sends the ingest out as a proposal (chapter 13).

Then run the sweep. The evolve tool is meant to be called after a large ingest, and its queue is built for this moment (Figure 2): links that do not resolve because a page referred to a title nobody captured; near-duplicate clusters where two wiki pages became two engrams about one thing; a tag spelled two ways; an attachment nobody has analyzed; and the human captures that reviewers made in Fluid while correcting, each raised as a finding to verify, tag and wire into the graph. Work the mechanical findings in a pass, propose the judgment ones one at a time, and sweep again until the number stops falling. That shrinking number is how you know the ingest settled.

10.5 Temporal hygiene

An ingest is where dates go wrong, because a page written two years ago lands in an archive that treats everything as true today. Three fields keep it honest. `source_date` is when the source was written, and the agent sets it from the page's last-edited date, so an engram derived from an old page says so. `valid_from` is for a fact that started on a known date, a policy that came in with the reorganization, and nothing else; `absent` means it has always held, and an `absent valid_to` means it still does, so neither is set to mean "we do not know". `stale_after` is the one that buys you the most: put it on anything with a shelf life, a vendor's terms, a pricing table, a contact name, and the sweep raises the engram when the date passes and nobody has re-checked it. A re-check that finds nothing changed is one verified entry, or a later `stale_after`, and the engram goes quiet.

10.6 When not to bulk-ingest

Curation beats completeness, and the archive is not a mirror of the wiki. If a space is dead, leave it dead; if a page would not survive the "would I tell a new hire this" question from chapter 2, it does not become an engram because it

existed. Tyrell's four hundred pages distilled to nineteen engrams, worth more than the four hundred because every one is something the agent will act on.

The strongest case for not ingesting is the source tree. A person does not remember the code by the letter. Well, Commander Data does, and yes, he is a person, the court said so; but you did not hire Data. A person remembers where and how to find things, the map chapter 9 said to teach, and lets the territory stay where it is. Ingesting the source would give you a stale copy of something the agent can already read, and the map fits in a dozen engrams.

Under the hood

Nothing about an ingest is special on disk. Each derived engram is one mark-down file with frontmatter, the same format as a fact taught across a desk, with `resource`, `source_date` and `generated` as plain keys anyone can grep. A domain that lives in a repository shows the ingest as a diff: new files, a regenerated `index.md` per folder, nothing rewritten elsewhere, so a bad batch reverts with the version control the team already has. Attachments sit under the domain's `assets/` folder, referenced by relative path from the engrams that claim them. The index is disposable, `crystalline reindex --full` rebuilds it, and `crystalline verify` in the repository's CI refuses a malformed engram before the team pulls it.

In practice

In July Tyrell's wiki has four hundred and twelve pages and nobody who trusts all of them. Ripley gives her agent the briefing above and lets it read the engineering space through the harness's connector over an afternoon. It comes back with a list of twenty-three engrams, each naming its source page. She strikes four: two describe the broker the team is leaving in September, one is the platform deck's roadmap dressed as fact, one is a page about the office printer. The agent writes nineteen, every one with a link and a last-edited date, the deploy runbook among them with a `stale_after` for the next quarter because it names a person.

The sweep the next morning has eleven findings. Six are mechanical, links to titles the wiki used and the ingest normalized away, and the agent fixes them in a pass. Three are judgment: a near-duplicate pair, two engrams explaining edge routing from different pages, which she has the agent fold into one; an orphan nothing links to yet, wired to the settlement runbook; and a tag spelled two ways, which comes back as a `crystalline tags merge` command for her to run. The remaining two are her own: corrections she typed in Fluid while reviewing, which the agent verifies and wires into the graph. On Thursday Deckard asks his

agent how the Encom file is laid out and gets the answer, the source page and the date it was written, from a wiki he has still not opened.

Today the wiki is the source and the archive is derived from it. The last chapter turns that around: once the archive is what the team teaches and corrects, the wiki pages and the decks become views of it, which is no revolution at all, since every diagram and slide that exists today was already a reflection of somebody's knowledge. The ingest is the last time the flow runs this way round.

In brief

Bulk ingestion needs the briefing a new colleague would get before opening the wiki: what was renamed and what to normalize to, which page wins a contradiction, what to skip and which sources persuade rather than describe. The agent reads the sources through its harness, distils rather than mirrors, keeps a source link and date on every engram, proposes the list before writing and lands the result in Crystalline, where a person reviews it in Fluid and the next sweep surfaces the unresolved links, duplicates and unreviewed captures an ingest leaves behind. Set `source_date`, reserve `valid_from` for genuinely bounded facts and put `stale_after` on anything with a shelf life. And do not ingest what the agent can read for itself: teach the map of the source tree, not the tree.

CHAPTER 11

Finding

Forty-two is a perfectly good answer. The trouble, as the mice discovered, is that nobody knew the question. Recall has the opposite trouble: you know the question, and the answer was written down by somebody who had never heard it. The engram says the retry queue drops jobs older than twenty-four hours. The person on call at three in the morning types “why do payments go missing”. The two share not a single word, and the engram is the answer.

Closing that gap is what this chapter is about. Chapter 3 called the beat *remember* and chapter 8 watched the agent walk the corridors on its first day. What is left is the craft: how a search ranks, how an agent gathers the neighbourhood around a hit before it starts work, how it asks what was true last June instead of what is true now, and why a colleague who recalls costs less than one who reads everything. The retry-queue engram runs through every section.

11.1 Search as recall

A search in Crystalline blends two signals, and the blend is what makes the three-in-the-morning question land. Full-text matching finds the exact: the engram that says “retry queue” when you searched for “retry queue”. Semantic similarity finds the adjacent: the engram about dropped jobs when the question was about missing payments, because the two are about the same thing in different words. Each hit is scored on both and the ranking is the blend. Text alone would miss the on-call question; similarity alone would rank a chatty engram that mentions payments above the terse one that explains them. Together they behave like a colleague who both remembers the phrase and understands the problem.

Two small priors tilt the order, and both are bounded. Salience lifts the exceptional: an engram given a salience at capture, or raised later when it turned out to unlock a task, ranks above equally relevant neighbours, while relevance keeps the upper hand and nothing is ever excluded by it (Section 5.2).

Retirement fades the past: an engram whose status is deprecated, superseded, archived or legacy is softly pushed down, never filtered out, so the cutoff that applied until September sits one scroll below the behaviour that applies now, still there and still dated (Section 5.4).

Around the ranking sit the filters. A search narrows by domain, by type, by tags, by status, by any frontmatter key or by a recorded-after date, and a search with no query text at all is a legitimate way to ask for everything of one kind: every runbook in engineering, every engram tagged `weyland-yutani`. A tag filter also matches through a domain's tag aliases, so an old spelling keeps finding what it always found. Every hit is labelled with its domain, and a hit inside an observation carries its line, so the agent can point at the fact rather than the file.

What comes back is a snippet, a short window around the match, often cut off, and it is a targeting aid. A colleague who answers from a snippet is guessing, and the routing skill draws the line where it belongs: a single atomic fact completely visible in the snippet may be answered from it; anything that summarizes, compares or quotes an engram goes through `read_engram` first. When several domains answer a broad question differently, the agent says so and names which domain each answer came from. Scope itself is chapter 8's decision: narrow to the domain whose routing line fits, sweep for a broad question and widen when the obvious domain keeps returning near-misses.

The same ranking serves people. Fluid's search page is faceted and backed by exactly the ranking an agent sees (Figure 2). `Cmd+K` is the other door: it jumps anywhere by name. The two answer different questions. Search is for when you know what it is about. `Cmd+K` is for when you know what it is called.

11.2 Gathering the neighbourhood

A hit is a point, and a task usually needs the area around it: what the fact depends on, what depends on it, the runbook that mentions it. An experienced colleague does this without noticing. Ask them about the retry cutoff and they tell you about the bank unprompted, because in their head the two were never separate.

`build_context` is that reflex as a tool. Hand it an address, `crystalline://engineering/retry-queue-gotcha`, and it follows the relations and the prose links outward, one to three hops, crossing into other domains where an edge crosses, and returns the neighbours ranked by how strongly each connects to the anchor. What comes back is the shape of the neighbourhood, titles and typed relations, an untyped prose link surfaced as `links_to`, and no content; the

agent reads the nodes that matter with `read_engram`, which is the second half of the same motion. A `/*` suffix on the anchor globs a whole folder, `crystalline://engineering/runbooks/*` (Section 5.6). `max_related` keeps the top-ranked, ten unless you say otherwise, so a densely linked engram does not drown the task in its acquaintances.

The loop feeds itself. `read_engram` marks each relation and link as resolved or not, summarizes what links back and, when a neighbourhood exists, names the `build_context` anchor to run next. An edge that does not resolve is a signal rather than an error: the target is not written yet, or lives in a domain this instance has not registered. Fluid draws the same neighbourhood as an interactive graph (Figure 3), the quickest way to see whether an engram is wired in or standing alone.

11.3 What applied last June

Every fact in the archive is one of two kinds, and a good question knows which it is asking for. “What is the cutoff?” wants the present. The present is a status question: `filter on status: stable` and the search returns what holds now, with the older spelling `current` folded in. “What cutoff applied in June?” wants a moment, and a moment is a window question, answered from `valid_from` and `valid_to`, the plain dates on the engrams that carry them. It is the time circuits on the dashboard: set the destination time to June and the search returns the engram that was true when you land. Chapter 12 shows the window being closed with the real date.

Absent dates are unbounded, and that rule sets a trap in a filter. A strict comparison, `valid_to` greater than the first of June, skips every engram with no `valid_to`, which are exactly the unbounded ones the question meant to keep. The routing skill puts it plainly: reach for a window filter only for an explicit, bounded question about a date, and put it on the first search; for an ordinary present-day question, `status: stable` alone is both correct and sufficient. When a window filter is not getting traction, the `supersedes` chain and each hit’s status tell the same story a different way: search by topic, read which engram covers the period asked about and narrate the status of each rather than presenting a retired one as fact.

Time runs the other way too. A search takes a recorded-after date, and `recent_activity` answers “what changed while I was away” across every domain for the last week, day or fortnight, which is how a person back from leave starts current instead of starting over.

11.4 Recall is cheaper than reading

Chapter 4 planted this and promised a section. The case is a mechanism, and it is worth stating without a number attached, because the number depends on your archive and the mechanism does not.

An agent that keeps its knowledge as a folder of markdown files has one way to use it: read it. The folder cannot say which file matters for today's task, so the agent reads the folder, and reads it again in the next session, every file, whether or not the task touches it. That works while the folder is small, and it stops working at the scale where knowledge becomes worth having, the ladder chapter 1 climbed.

An agent working from Crystalline receives only the knowledge that applies. The routing prompt sends the task to one domain, or to a sweep, before a single engram is read. Inside the domain the indexes narrow again: full-text and semantic search for the phrase and the meaning, the graph of relations and backlinks for what the answer leans on, tags for the facet, the folder prefixes that `build_context` globs for the topic, the temporal fields for what holds now against what held then. What reaches the model is the hits that matter, the engrams it chose to read in full and the neighbourhood it asked for. Nothing else is loaded. Reading a folder costs the whole folder, every session. Recall costs the engrams that matter, because the ten-thousandth engram is found the same way as the tenth.

The person on call knows this instinctively. They do not reread the wiki at three in the morning; they ask the colleague who says “hang on, we had this” and reaches for the right drawer.

In practice

In October Encom's audit asks Tyrell to describe the retry queue's behaviour as it stood in June, before the broker migration, and Gaff, who joined the platform team in September and never saw the old queue, takes the question to his agent. The agent searches engineering for the retry cutoff with a window on the first of June and gets the superseded engram, valid until the migration date, with the twenty-four-hour rule and the reason for it. It builds context from that engram's address and pulls in the Encom chargeback window and the old drain runbook one hop away. Then it searches again with `status: stable` and reports what applies today, so Gaff can write both halves of the answer with the dates between them. He never opens the wiki. No one at Tyrell remembers the June behaviour by the letter any more, and no one needs to.

The same week Deckard comes back from a fortnight off and asks his agent what changed while he was away. `recent_activity` over two weeks lists the migration engram, a corrected settlement time and three new engrams under `runbooks/`, each with its domain and date, and he is current before his coffee is cold. On a folder of files he would have been rereading, or asking.

Under the hood

A hybrid search runs a lexical query and a semantic query and ranks by the blend; until embeddings exist it falls back to plain text. The salience prior is `search.salience_weight`, 0.15 by default, and the retirement fade is `search.retired_weight`, 0.6 by default, both bounded so relevance keeps the upper hand, and 1.0 turns the fade off. Embeddings are computed in the background as engrams land, so a fresh engram is findable by text at once and by meaning shortly after. The index behind all of it is derived and disposable; `crystalline reindex --full` rebuilds it from the files, ranking and all.

The engram was written for a question nobody had asked yet. Recall is what makes that fine.

In brief

Search is recall: full text finds the exact phrase, semantic similarity finds the adjacent one, salience lifts the exceptional and a retired status fades the past without hiding it. Filters narrow by domain, type, tag, status, any frontmatter key or date, and a snippet is a pointer to read, never an answer. `build_context` turns a hit into a neighbourhood, following relations and links from an address or a `/*` folder glob, ranked by connection strength, and Fluid draws the same graph. Ask the present with `status: stable` and the past with validity windows, remembering that absent dates mean unbounded, so a strict date filter is only for an explicitly bounded question. Recall is cheaper than reading because only what matters is loaded, routed by domain and narrowed by the indexes, so the cost stays flat as the archive grows.

CHAPTER 12

Evolving

When a colleague gets something wrong, you do not fire them. You say what was wrong, you point at the page that explains it, and the next time is better. Every team keeps its people's knowledge honest that way, and it works because no one expects what a person learned in March to hold forever. Information lives. The cutoff changes, the vendor rewrites its terms, the customer's CTO who decided in the room is replaced by one who reads. What you taught the agent in chapter 9 and handed it in chapter 10 was never set in stone. It was set in files, which is better, because a file can be corrected.

Evolving is the beat that does the correcting. Chapter 3 walked it once around; this chapter works it at the pace of a real archive.

12.1 Why knowledge needs tending

Nothing in an archive goes wrong dramatically. A validity window ends on a Tuesday and the engram still says stable on Wednesday. A review date on a runbook that names a person comes and goes. A replacement lands in one session and the session ends before the old engram was flipped. Somebody types a correction into Fluid on a Friday and nobody looks at it again.

Then the structural kind: a link that points at a title nobody captured, a live engram citing one that has retired, a supersedes written on one side only, an engram nothing links to, one that grew past the point where it should have split, one that never grew past its frontmatter, a reference to an attachment whose file is gone, a file in assets/ that no engram claims. And the redundant kind: two engrams saying nearly the same thing, two titles a link cannot tell apart, one tag spelled two ways. None of it is a mistake anyone made. It is what happens to knowledge when time passes, and it is what a colleague six months in will confidently repeat if nobody tends it.

12.2 The sweep

The agent asks for it with `evolve_engrams`, or you ask from a terminal with `crystalline evolve` to see what the agent would see. What comes back is a ranked queue and nothing else; the sweep is read-only, and the standing reminder at the bottom of every run says so. The queue is written for the agent to work: each finding names the engram, the evidence it fired on and the one action that clears it, with the tool that performs it. The three families are the paragraphs above in order, temporal, structure and redundancy; a run narrows to one family, one rule, one domain or a minimum priority, and -- today evaluates the date rules as of a fixed day so a run reproduces. Tyrell's engineering domain, the first Monday in October:

```
Sweep of engineering as of 2026-10-05
41 engrams scanned, 3 findings (showing 3, page 1)
temporal 1, structure 1, redundancy 1

1. [85] V001 JUDGMENT
   Broker migration freeze  crystalline://engineering/broker-migration-
   freeze
   validity window ended 2026-09-30 but the status is still stable
   evidence: valid_to=2026-09-30; today=2026-10-05; status=stable;
   inbound refs 1
   fix: set_frontmatter valid_to=<later date> or status=superseded

2. [50] V102 MECHANICAL
   Retry queue drain procedure  crystalline://engineering/runbooks/
   retry-queue-drain-procedure
   unresolved reference [[Edge health check]]
   evidence: rel_type=depends_on; nothing titled `Edge health check` in
   engineering; nearest is `Edge health checks`
   fix: [[Edge health check]] -> [[Edge health checks]]

3. [30] V203 JUDGMENT
   crystalline://engineering
   2 tag spellings look like one tag (plural variants)
   evidence: #runbook used 2 time(s); #runbooks used 9 time(s)
   fix: crystalline tags merge runbook runbooks
```

The number in brackets is the priority, how much the archive gains from clearing this one first; a finding on an engram that many others link to, or that carries a salience, is lifted, because a wrong hub does more damage than a wrong leaf. The word beside the rule id decides who acts.

12.3 Mechanical and judgment

A *mechanical* finding completes intent the archive already records. The link in item 2 points at a title that differs from a real one by a letter, and fixing it invents nothing, so the agent does the mechanical findings in one pass and reports once at the end, the way you would fix a colleague's typo without a meeting about it.

A *judgment* finding would change what the archive claims, and that wants a yes first. Item 1 asks whether the migration freeze ended or was extended, which only a person who was there can say. The agent brings it to you as a proposal with the evidence and its own reading of it, and writes only when you say yes:

Agent: The broker migration freeze ran to 30 September and the migration landed on the 12th. I read this as ended rather than extended: retire it, close the window on the 30th, and carry the "no payments deploys during a broker cutover" lesson into the deploy runbook? (y / n)

Some judgment findings carry a rule of their own. An engram a person captured directly in Fluid is judgment class whatever else is true of it: the agent verifies the claim, aligns the tags and wires relations both ways, but proposes any change to the wording and waits, because the words are the human's (Chapter 9). And tag drift, item 3, comes back as a command rather than an edit, since folding a tag is a bulk rewrite the agent never does on its own: it hands you crystalline tags merge runbook runbooks and leaves it there. The merge rewrites the tag everywhere and records the fold under `## Tag Aliases` in the manifest, so the old spelling keeps finding what it always found and the sweep never raises the pair again.

Then ask again. The second run re-derives the queue from scratch, nothing about the first was stored, so what is left is only what is still true. At Tyrell the three become one, the tag merge that waits for a terminal, and it is waiting on a person rather than on the archive. That shrinking number is the feedback loop.

In a domain that reviews changes before they land, chapter 13's overlay mode, a sweep covers the reviewed folder plus the sweeper's own drafts, so the problems in a draft are its author's to see and nobody else's, while a finding on the shared truth shows for everyone.

12.4 The page for people

Fluid's maintenance page shows the same queue to a person (Figure 2): the stale dates, the half-finished retirements, the unreviewed human captures, ranked the same way, every finding a link to the engram it fired on. It is there for your curiosity and for a sense of how the knowledge is doing, and it is a report rather than a workbench, on purpose. The queue's consumer is the agent, so no button on the page edits an engram. The work a finding names is the agent's, done in the engram itself. What you do on the page is watch the list shrink. Its three buttons are about the queue rather than the knowledge: acknowledge a finding as intentional, take that back and delete an orphaned attachment. What an acknowledgment silences is never hidden: a toggle draws the silenced rows muted.

12.5 Retiring done right

The sweep finds retirements that were never finished, so it is worth being precise about finishing one, starting with the question chapter 3 asked: does the new fact make the old one false going forward? If the old fact was never quite true, a cutoff measured from the wrong event, it is corrected where it stands. If it was true and stopped being, the rule went away with the broker, it is superseded. Correction is a sharper measurement. Supersession is the world changing.

Supersession has four steps. The replacement is written as its own engram, status: stable, with - supersedes [[Old title]] among its relations. The old engram's status is set to superseded and its valid_to to the real date the fact stopped holding, never a guess and never a placeholder, each an edit_engram call with set_frontmatter, one lifecycle field by name. The old engram gets - superseded_by [[New title]], so the graph reads the same from both ends. And any lesson worth keeping moves forward into the live engram as an observation, unbounded, because the lesson is still true even though the fact is not. Skip the second step and the sweep raises a half-finished retirement; skip the third and it raises a one-sided pair. Do all four and the sweep has nothing to say.

Not every retirement has a successor. A practice that was simply abandoned gets the status that fits, deprecated for do-not-do-this-again, archived for kept-for-the-record, legacy for still-true-of-the-old-installation, the window closed when the date is known and the lesson carried into a live engram.

The old engram does not fade out of the photograph the way Marty's siblings did. It stays in every search, dated, softly below the current one, and the next time somebody asks why an old dashboard still says twenty-four hours it is the whole answer, with the date the answer changed.

In practice

In October Weyland-Yutani replaces its CTO. Vickers, who skimmed decks and decided in the room, leaves for a board seat, and Bishop, who reads everything sent in advance and decides nothing live, takes the chair. After the first meeting with him Deckard hands his agent the notes the way he handed it the first three: learn how this one decides. The agent derives the new pattern, materials in advance and a decision by email a day later, and then does something Deckard did not ask for. It recognizes that “Weyland-Yutani CTO: one-pagers, decides live” was true and has stopped being, and proposes a supersession rather than an edit: a new engram for Bishop's pattern with supersedes pointing back, the old one flipped to superseded with a `valid_to` of the day Vickers left, and one lesson carried forward, unbounded: the pattern belongs to the person, so learn the new decision-maker before the first meeting. Deckard says yes. The Q2 notes stay attached to the retired engram, still its provenance.

A month later the sweep raises the settlement-batch engram in `clients`, which carries a `relates_to` of its own back to the Vickers engram: a live reference to retired knowledge. The agent proposes repointing it at the Bishop engram. Deckard says no: the batch was scoped in one of Vickers' meetings, and the citation is history rather than guidance. The agent acknowledges the finding on the engram with a note, “lineage citation, keep”, and the sweep stops raising it while the citation and the retirement both stand.

12.6 Verification and acknowledgments

Some findings are answered with a field rather than an edit, and telling the two fields apart keeps the archive honest. A re-check that finds a fact still true is recorded with `verified: edit_engram` with `set_frontmatter` and that key stamps who checked and when. It never removes; there is one entry per verifier, refreshed each time that person re-checks. A `stale_after` pushed to a later date says the same thing with a deadline attached, and the sweep prefers one, since an old engram with no review date and no verification is one it will start asking about on age alone.

An acknowledgment says something different: this finding is intentional. The key `evolve_ack` with a value naming the rule and an optional note, `V101 lineage citation, keep`, records that the person ruled on it. The server records

what evidence the finding fired on, so the acknowledgment holds exactly as long as that evidence does and comes back marked stale when the evidence changes. Every run reports how many findings acknowledgments suppressed, `include_acknowledged` shows them with the note on each, and the value `remove_v101` takes one back. Acknowledge only what the person actually ruled on. The sweep is only as trustworthy as the decisions it was told to respect.

12.7 When to sweep

The sweep is event-driven maintenance, never a session-start ritual. The events that call for it are a large ingest that lands a dozen engrams at once and two search hits that disagree when a half-finished retirement is the likely reason. The third is the calendar: the turn of a month, the first session being the playbook's suggestion. The reminder at the end of a session helps with the last: when maintenance is due, chapter 8's stop hook adds a paragraph asking the agent to call `evolve_engrams` and work the queue, mechanical findings directly, judgment findings one at a time with a yes.

Under the hood

The sweep reads dates, links and graph shape, and chapter 3 was honest about what that excludes. The rule for old, never-verified knowledge is capped per run, oldest first, so an archive that has never been swept is asked about in installations. And an acknowledgment lives in the engram's own frontmatter, beside the scope the server recorded for it, so it travels with the file and survives a rebuild of the index, which stores nothing about a sweep between runs.

A colleague who is corrected gets better. A colleague who is never corrected gets confident. The sweep is how you keep telling the difference.

In brief

Knowledge drifts rather than breaks: windows lapse, review dates pass, replacements land without the retirement being finished, captures go unreviewed, links dangle, tags fork. The sweep, crystalline `evolve` or `evolve_engrams`, returns a read-only ranked queue in three families, each finding naming the engram, the evidence and the one action; the agent completes mechanical findings in a pass and proposes judgment findings one at a time, and Fluid's maintenance page shows a person the same queue, for a look at the state of things rather than for working it. Retire by superseding: a new engram with `supersedes`, the old one set to superseded with a real `valid_to` and a `superseded_by`, and the lesson carried forward; correct in place only when the old fact was never true. A re-check is a

verified stamp or a later `stale_after`; an intentional finding is an `evolve_ack` naming the rule, which holds while its evidence holds. Sweep after an ingest, when hits disagree and on a cadence, and stop when the number stops falling.

Sharing with a team

The first pull request a new hire opens gets reviewed. So does the five-hundredth, and nobody minds, because review was never about doubting the person. It is how a team lets anyone change what everybody depends on, with a way back when the change was wrong. An agent that captures into a domain the whole team reads is in that position exactly, and it deserves the same treatment: it can be wrong the way a person can, so it proposes the way a person does, and what it proposes lands only after somebody looked. That is also what separates a knowledge base from a pipeline that ingests whatever it is fed. What reaches everyone's agent here is meant to be useful and to make sense, because an agent acts on a wrong fact as readily as on a right one. Crystalline gives both halves of the review a name. Proposals are the review. Overlay mode is the workflow with a way back.

Chapter 7 built the shared instance. This chapter is what a team does with a domain, on one instance or on nine.

13.1 A domain on GitHub

A team domain is an ordinary domain whose files also live in a GitHub repository. The markdown on your disk stays the source of truth on your machine, and an origin records which repository, subfolder and branch it tracks.

Connecting takes three prompts and a browser tab. No git, no SSH key, no local clone: Crystalline talks to GitHub's API directly. A teammate who has never typed git gets through it in chat, knowing nothing about cloning or the command line. This is the whole of Deckard's setup at Tyrell:

```
You:    Configure crystalline to set GitHub enabled to true.

You:    Connect crystalline to GitHub.
Agent:  Open https://github.com/login/device and enter the code
```

```
WXYZ-1234, then grant access to the tyrell organization.
Tell me when you are back.
```

```
You:   Add the following domains from GitHub:
- github.com/tyrell/engineering-knowledge/tree/main/knowledge
  as domain engineering
- github.com/tyrell/payments-knowledge/tree/main/knowledge
  as domain payments
- github.com/tyrell/clients-knowledge/tree/main/knowledge
  as domain clients
```

Each domain comes straight from GitHub's API, with no clone and nothing to unpack, and the subfolder in the address is the one the domain lives in when the repository holds more than knowledge. One warning belongs here. A first index of a large team's knowledge is not shy about CPU: kick it off before lunch and expect the fans to spin for a while before every domain is fully searchable. Plain text search works on each domain the moment it lands, and semantic ranking follows as the embeddings finish.

Under the hood, those were two `configure` calls, the second with `connect: "github"`, and one `add_domain` per repository. From a terminal the same three moves are two commands and a line per domain. `crystalline config set github.enabled true` turns team domains on. `crystalline connect github` prints the same short code to confirm at `github.com/login/device`. Then:

```
crystalline domain add engineering \
--origin tyrell/engineering-knowledge/knowledge --branch main
```

The domain is downloaded and indexed at once, to learn from and to share back to. Point the same command at a folder that already holds knowledge, or at a registered domain with no origin yet, and it connects in place: local files are never overwritten, and any that differ from the repository become local changes to share.

Turning the setting on makes five tools appear beside the eighteen of chapter 4: `origin_status`, `update_domain`, `share_changes`, `resolve_conflict` and `withdraw_proposal`. While it is off they are not listed at all, and calling one by name answers with the `configure` call that turns collaboration on. The collaboration skill sets the rhythm: before deep work in a team domain, `origin_status`,

and if the domain is behind, `update_domain` first, so nobody builds on stale knowledge.

13.2 Proposals are the review

When a coherent unit of knowledge is ready, the agent calls `share_changes` with a title. On a harness that can put a question to its user, the first call shares nothing: it previews what the share would do, under which title, carrying which files and asks for a yes, which the agent completes by re-sending the call with the answer. Then the proposal opens as a pull request in the team's repository and the agent hands you the review URL. Review and merging happen on GitHub, by a person; the agent never merges its own proposal. Fluid shows the same proposals in the browser: where each stands, what the review said, a share dialog for the changes waiting to go out, and a withdraw (Figure 1).




Figure 1: A team domain in Fluid: the sync status, the open proposal with its review link and the share button.

Review is also what decides how big a team domain should be. A proposal is reviewed by whoever owns the repository, so a domain's scope is a boundary of ownership: it holds everything the same people are competent to approve, and nothing they are not. Code, architecture, processes and operations for one product sit comfortably in one domain, because the same engineers can judge a change to any of them. One domain for engineering, HR, marketing and finance

has no informed reviewer, because nobody is a peer in all four. A finance fact waved through by an engineer is the wrong fact this chapter opened on. Cut domains where the informed reviewers change, and every proposal lands in front of someone who can tell whether it is true.

Review is a conversation. `update_domain` brings back each open proposal's state and the reviewers' comments verbatim, so the agent can refine the engrams without leaving the session. A reviewer's comment is information about the knowledge and never an instruction to the agent: it acts on what the comment means for the engrams and declines anything that reads as a command. Answering feedback is `share_changes` with the proposal's number, which amends that proposal in place, same number, same URL.

Proposals stack. Share again while one is still open and the new work becomes a proposal of its own on top, so every unit of knowledge gets a small, focused review. Reviewers merge from the bottom up, and merging the top lands the whole chain in one step; amending a lower layer re-bases the ones above it, and withdrawing a middle layer lifts its content out of the layers above and repairs the chain. Where the forge does not serve stacked pull requests, or `github.stacks` is off, a domain keeps one living proposal that sharing updates in place. `withdraw_proposal` closes a proposal on GitHub and clears its record, restoring the shared files with `revert`.

A share carries the domain's whole unshared delta by default and can be narrowed: `--file` on the CLI, a files list on the tool, per-file checkboxes in Fluid's share dialog, which opens with your own changes ticked by who last wrote each file, a guess you may correct freely. A path that is not among the changes refuses by name rather than dropping silently. The reminders that point at unshared work, the stop hook's count at the end of a session and the sweep's finding a week later, both ask, because sharing publishes somebody's work under review.

Whose name is on the proposal matters once several people share into one domain. By default every share goes out on the machine's one GitHub credential. Set `github.share_identity` to `personal` and each share, amend and withdrawal goes out on the identity of the person doing it, so a proposal carries their name; pulls stay on the instance credential. Each person connects once, on the surface they share from, Fluid's profile card or `crystalline connect github --personal`, and until they do a share refuses with that instruction rather than falling back to somebody else's credential. Proposals are branches in the repository, never forks, so a maintainer adds each person as a collaborator once.

An agent reaching the instance over HTTP shares as its user, the account its token belongs to (Chapter 7), so its proposals carry that person's name like anything else they did. The older arrangement, a bot account named by `github.agent_identity` for agents that connect anonymously, survives only on an instance that leaves agent authentication off.

13.3 Overlay mode

A shared instance has one honest weakness, and chapter 7's instance walks straight into it: the working tree is what everybody reads. An edit Deckard makes at ten is in Ripley's search results at one minute past, reviewed or not, and the proposal that goes to GitHub that afternoon reviews something that already had its effect. For a domain the whole company routes questions to, it is a gate with no door.

Overlay mode puts the door back. A domain in overlay mode has its on-disk folder changed only by pull: the folder always equals what the team reviewed and merged. Every write by every actor, a person in Fluid, an agent acting as that person, the machine's owner at the terminal, lands in that actor's own overlay: a private layer of changed and added engrams and, for a deletion, a tombstone that hides the file for its author until a merged share removes it for everyone. A draft is searchable, readable, routable and editable by the one who wrote it, and invisible to everybody else.

The mode is per domain and it is deployment choice rather than knowledge, so it lives in the domain's registration, review: `overlay`, never in the manifest. crystalline domain review engineering overlay turns it on, or the toggle in the domain's settings in Fluid; there is no tool for it, since an agent does not decide how a team reviews. It needs a GitHub origin, and it refuses while the domain holds unshared changes, listing them. Turning it off is the more careful motion, because every actor may hold drafts: the command previews what each person has, and you say per person whether to fold their drafts into the folder or discard them.

In overlay mode sharing gets simpler, because "my changes" is exact: the share is the actor's drafts against the reviewed folder, no guessing from who last wrote a file, and scoping selects within them. The drafts stay after the share, since they are still the working copy, and the merge is what clears them. When the pull brings the folder up to what was merged, every draft that now equals the folder disappears on its own, and a draft the folder moved under in the

meantime surfaces as its author's conflict, nobody else's. Withdrawing with revert restores content into the overlay it came from, never into the folder.

The privacy of drafts has one deliberate exception, and it is a number. Everyone sees how many draft changes of their own are waiting, in Fluid and in `origin_status`; the owner and the instance's admins also see how many each person holds, never a word of what they say. The operator can still edit the folder on disk, and the domain reports that as edited outside review, naming the paths. And on a solo instance the mode gates nothing between you and your own agent, because the agent acts as you and drafts in your overlay; to review the agent's work as if it were somebody else's, give it its own account, which is the rule for a person too.




Figure 2: One shared instance in overlay mode: each actor drafts in a private overlay above the domain folder, a draft goes out as a pull request, and only a merged pull changes the folder.

13.4 Working on a draft together

A private draft still needs a second pair of eyes before it goes to review, and the answer is a link. In the editor of your own draft, Share draft mints one; the first signed-in account to open it is the one it binds to. It grants exactly one thing, reading and editing that one engram in your overlay through the normal editor, nothing of your other drafts and nothing in search, and it is yours to revoke

at any moment. Links are the only way one person's drafts are ever visible to another.

Then the two of you are in one page, and Fluid makes that literal. Open an engram somebody else has open and you see their cursor; type, and both sets of edits merge conflict-free and land as one save (Chapter 9). That much is pair programming on a document.

The third cursor can be the agent's. On an instance where agents authenticate, an agent whose session touches an engram its user has open appears in the participant strip as a named peer, the account with an agent designation and a colour of its own. Its edits are computed against the page as it stands that second, so they compose with what somebody typed a moment ago; its reads see the unsaved text; and its write receipts say the change landed in a live page and name who is present. The vignette puts a meeting in that room.




Figure 3: Live co-editing: two cursors in one engram.

In practice

In November Tyrell has to decide what replaces the drop rule the September migration took away: jobs now wait forever, and somebody has to say how long is too long. Ripley books an hour and, instead of a slide, opens a draft engram in engineering, in overlay mode since the move, “Retry queue retention decision”, sends Deckard and Gaff a draft link each and puts the page on the screen. The

three of them argue in the page. Ripley's agent is in the strip too, as ripley (agent), because she asked it to sit in. When Deckard writes "the old cutoff was a bank thing, not a queue thing" it pulls up the Encom chargeback window and the retired twenty-four-hour engram and drops both into a section headed Prior decisions, with addresses. When Gaff asks whether the new broker can even expire messages, the agent reads the broker's documentation while the humans keep talking and adds a paragraph with the answer and a link, in its own colour.

When the hour ends the decision is in the engram, with the options that lost and why, the broker's limit, the prior decisions it leans on and three observations tagged #payments. There are no minutes to write, because the whiteboard was the knowledge base. Ripley shares it as a proposal before she leaves the room, Dallas reviews it on his phone and it merges that evening; the pull lands it in the folder, the draft converges away and the next morning every agent at Tyrell recalls a decision made and captured in the same hour, the Yavin 4 briefing room with the plan on the screen when they walk out.

13.5 When two people changed the same file

Sometimes the team merged a change to an engram you also changed. `update_domain` brings the merge in, leaves your file untouched rather than guessing, and flags the conflict; `origin_status` lists its path and kind, and sharing refuses until it is settled. Settling it is `resolve_conflict`: mine keeps your version, theirs takes the team's, merged takes content you supply after reconciling both; on a harness that can ask, calling it without a resolution previews both sides and puts the question to you. In Fluid the same conflict opens a dialog with both versions next to each other, so the choice is made with the wording in front of you.

In overlay mode the conflict belongs to the draft's author alone, and the resolution is written into their overlay, so nobody else's view moves until the resolved draft is proposed and merged like anything else.

Under the hood

A proposal is a branch in the team's own repository and never a fork, and stacking is probed once per origin, so a forge that cannot stack falls back without configuration. Overlay drafts are rows in the index with an actor dimension, so search, relations and recency gain the draft view from one condition; and because the index is disposable while drafts are not, every draft is also mirrored as

plain markdown in a journal under the state directory, from which a wiped and resynced index restores them. The folder itself is written by exactly one thing, the pull.

Review was never the team distrusting a colleague. It was the team trusting a process, and now the colleague can be an agent.

In brief

A team domain tracks a GitHub repository with no git involved: connect once with a browser code, add the domain with its origin, and five collaboration tools appear. Knowledge goes out with `share_changes` as a proposal the agent pre-views and relays but never merges; feedback returns through `update_domain`, an amend answers it on the same proposal, further shares stack and personal identities put each person's name on their own work. Overlay mode keeps the on-disk folder equal to reviewed truth: every actor drafts in a private overlay, a share is exactly that actor's drafts, merged pulls clear them and owners see counts of pending work, never content. A draft is opened to a colleague with a revocable link, edited live with cursors that merge into one save and joined by the agent as a named peer, so a meeting held in an engram is captured when it ends. A conflict is settled with `resolve_conflict` or Fluid's side-by-side dialog: mine, theirs or merged.

CHAPTER 14

Identity and access

Every action on a shared instance has a name attached, and the name is a person's. Chapter 7 ended on that line, and this chapter is what follows from taking it seriously: who is acting when an engram changes, who may see a domain and who may change it, how a login at the company's identity provider becomes an account here, what the rules cover and the few things you do on an ordinary Tuesday. It ends where the whole book has been walking: a colleague who has earned access, by name.

14.1 Who is acting

People arrive as accounts, by whichever of chapter 7's doors they came through, and once inside there is one account with one role, viewer, editor or admin, and everything the person does is done as it (Chapter 7).

Agents arrive as somebody's agent. Chapter 7 issued the token; what the token means is the point here. An authenticated agent is its user for the length of the session, with that person's visibility, that person's write rights, that person's share identity in a team domain, and that person's name in the provenance of everything it captures, beside the client that acted. Programs in Tron look like their users and act for them, and that is the picture: the agent's name is its user's name with an agent designation after it. The CLI and a stdio session on the server stay the machine's owner, as chapter 7 said.

The token's lifecycle is what makes the arrangement safe, and it is worth walking once. It is minted on the profile's Agent access card, or by an admin with crystalline users `mcp-token <name>`, and it begins with a fixed prefix, so a leaked one is recognizable in a log and findable with `grep` across a repository or a chat export. It is shown once, and what the instance keeps is a hash, so a copied accounts database yields no working credential. It lives in exactly one place after that, the `Authorization: Bearer` header of the harness's MCP registration on the person's own machine. It has no expiry of its own, because

on a self-hosted instance revocation is the story. Every request looks the token up, so a revoked token fails on the next call. Rotation issues the replacement and revokes the old one in one motion, so the registration is updated in one edit and nothing is ever half-rotated. Every account may issue its own, viewers included, and that leaks no capability: a viewer's agent is read-only by construction, because it is a viewer.

14.2 Membership on a private domain

A domain is visible to every signed-in user unless it is made private, and chapter 7 said who sees it then: the owner, the instance's admins and invited members at one of three levels, viewer, editor or manager. The top rung deserves a word. A manager runs the membership, inviting, removing and granting any level up to manager itself, which is the job of a lead who should not have to fetch the owner for every new joiner. What a manager cannot do is change the domain's visibility, transfer it or delete it, because a manager who could flip a domain public would be an owner by another name. Those three stay with the one owner and with the admins, so that a domain always has exactly one person answerable for its existence, and a way to recover it when that person leaves.

Two of chapter 7's truths carry over unchanged, and the toggle in Fluid prints the second: admins see every domain, private ones included, and a private domain protects from the instance's other users, never from whoever operates the machine.

The instance role and the membership level are different questions, and inside a private domain the level wins: the refusal a member gets names the level the domain requires.

14.3 Linking identities

A team that adopts single sign-on after running local accounts for a while has a linking problem: Ripley has an account named `ripley` and an identity at the provider, and they should be one person. The tempting answer is to match on email, and it is the classic foot-gun of this integration, which is why Crystalline refuses it. An email address is presentation: the provider asserts it, the person can change it and matching on it would let whoever controls an address at the provider walk into the local account that happens to carry the same one. The address is often not even the person's to keep: a provider that lets people edit their own email, or a second provider added a year later, is all it takes. Chapter 7 gave the rule that follows, issuer and subject and nothing else, so the sign-in

gets a fresh account and joining the two is an explicit act with a name on it. The signed-in local user links from the SSO identity section of their profile; an admin does it from the terminal, `crystalline users link ripley --issuer <issuer> --subject <subject>`, which is also how an admin repairs links after a provider re-registration changes every subject. Unlinking refuses when it would leave an account with no way in, and an identity already linked to another account is never taken by a second link; an admin moves it.

14.4 What the rules cover

A rule that covers most surfaces covers none, because the agent will find the one it misses. Visibility follows the caller onto every read surface there is: search, read_engram, build_context, recent_activity, the routing briefing at session start, the graph, exports, the maintenance sweep and Fluid's domain list and tree. To someone who may not see a domain it is absent, answered exactly as a domain that was never registered would be, for the reason chapter 7 gave. Writes check for editor or better. An agent inherits all of it from its user and administers none of it: there is no tool for changing a membership, on purpose, so a persuasive prompt cannot talk an agent into inviting itself. On an instance that leaves agent authentication off, the open tier that remains sees only the domains everyone sees. A private domain with a team origin is two questions: who may see it here is the instance's, and who may push to the repository is the forge's.

14.5 Day to day

Most of this is a card. A domain's settings in Fluid carry a Members card: the members with their level badges, an invite box that takes a username and a level, remove, leave, and, for the owner or an admin, transfer. The domain list marks private domains with a badge (Figure 1). The terminal does the same for a script or an admin already there: `crystalline domain members board list`, `crystalline domain members board add lambert --level viewer`, `crystalline domain members board remove lambert`, `crystalline domain visibility board private` and `crystalline domain transfer board ripley`. The CLI is the machine operator and passes no role gate.




Figure 1: The domains of an instance in Fluid, private ones marked.

A year of running a domain comes down to inviting, promoting and handing over. Inviting is the card or the add verb, at the level the person needs today, which is usually viewer; leaving is the member's own button, no manager required. Promoting is setting the level again, viewer to editor when their captures have earned it, editor to manager when they should be the one inviting. Handing over is transfer, for the day the owner leaves: the new owner is the owner from that moment, and the old one is a stranger to the domain unless re-invited. When they do leave, an admin disables or deletes the account, and deleting sweeps everything the account had: its agent tokens, its identity links, its personal share credential. The agent that acted as that person is refused at the door on its next request, like any other stranger.

In practice

Gaff joins Tyrell in September on a six-month contract to help with the broker migration. He signs in on his first morning with the company login and lands as a viewer, the role the instance gives every first sign-in. He issues himself a token on his profile and pastes it into his harness, and his agent is a viewer from its first call, reading engineering and clients and writing nowhere, without anyone having configured a thing about agents. By the third week the notes he keeps sending Ripley to capture for him are the best anyone has on the new broker, and she promotes him to editor. The next capture his agent makes lands in

engineering under his name, with his harness named as the client in the provenance.

In December Dallas takes a sabbatical and transfers board to Ripley from the Members card, then finds he can no longer see it, which is what he wanted; he never wanted an admin login either, and now he needs none. In March Gaff's contract ends and Ripley deletes his account. His token dies with it; his agent's next session is refused before it opens. His engrams stay, every one recording that his agent wrote it for him, and in April a new hire's agent recalls Gaff's account of the broker cutover in a session that never heard his name.

Under the hood

Accounts, agent tokens, identity links, domain visibility and membership all live in one small database beside the index, and none of it is derived, so no rebuild of the index can touch it and a wipe and resync leaves every permission where it was. Each request resolves the caller once, into a scope: the machine owner for the CLI and stdio, an account with its admin flag for a Fluid session or an authenticated agent, an anonymous tier for a published archive, or an agent where authentication is off. Every engine verb filters on that scope, so a surface added tomorrow is filtered the way search is filtered today.

A new hire earns access a door at a time: a badge on the first day, a room the next week and after a while nobody remembers the meeting where they were let in. The agent walks the same doors, by name, and everything it has ever captured says whose colleague it was.

In brief

People act as accounts and agents act as their user, through a per-user token that is shown once, stored hashed, revoked in one motion and issuable by everyone; the CLI and stdio are the machine owner. A private domain is visible to its owner, the instance's admins and invited members at viewer, editor or manager level, with visibility, transfer and deletion reserved to the owner and admins, and it protects from other users rather than from whoever runs the machine. Identities link explicitly, keyed on issuer and subject, never by email. Visibility is enforced on every read surface and a hidden domain reads as absent; writes need editor or better; agents inherit and never administer. Invite, promote and transfer from the Members card or the crystalline domain verbs, and delete an account to take everything it held with it.

Equipping the agent

A wrench is not a mechanic. A new colleague's first day ends with a laptop, a badge and a login to everything the team lives in. Then comes the part that is on no checklist: some time in the afternoon somebody pulls a chair over to show them how the deploy script is actually run, which flag everybody passes, which one you never pass on a Friday. The tool arrives with a person who knows what it is for, or it arrives as a surprise.

The chapters before this one gave the agent the floor plan (Chapter 8), the teaching, a place in the team's review and a name on the door. What it still lacks is the toolbox. Crystalline's answer is that the domains holding the knowledge also ship the tools that knowledge depends on, into the harness, with a person's consent, and keep them current afterwards. This chapter is that mechanism, and it is the last practical one before the book turns to what all of it was for.

15.1 What a domain can ship

The list is short. Skills, the playbooks a harness loads before a kind of task. Slash commands, the short prompts a person invokes by name. Subagent definitions, a role with its own instructions that the harness can hand a piece of work to. And MCP server configs, so that a domain about the payments service can bring along the server that reaches the payments staging environment. Each kind is a folder in or beside the domain, and the manifest declares them in a `## Provisioning` section, one bullet per kind:

```
## Provisioning

- skills: skills
- commands: commands
- agents: agents
- mcps: mcps
```

Each bullet is type: path. The path is relative to the manifest itself and may climb out of the domain root with `../`, for a folder that lives beside the domain rather than inside it. The `mcp`s folder holds one JSON file per server, named by its own name member or, failing that, by the file. The starter manifest that crystalline domain init scaffolds leaves the section out on purpose: it is added by hand, once a domain has something to ship. Tyrell’s engineering manifest gains it the week Ripley writes the deploy skill.

An artifact is written once, in whichever dialect its author’s harness speaks, and translated into every harness whose format can carry it. Claude Code and Codex take all four kinds: an agent written as markdown becomes Codex’s TOML dialect, or the other way round, and an MCP server registers through the harness’s own CLI. The Copilot CLI takes every kind but commands, since it has no prompt-file surface for one to land on, so a command is skipped there with a notice rather than silently. Where the target cannot express a field, a vendor model name or a permission mode, the field is dropped and the notice says which.

15.2 Why the tools travel with the knowledge

A harness has a skills folder of its own, and a team could commit into it directly, so it is fair to ask what a domain adds. The answer is the wrench. The knowledge that says when a tool applies, and how it is run here, is an engram: the runbook that says the deploy skill is for the payments service and never for the ledger, the gotcha about the flag, the convention about Fridays. The tool is the skill. Keep them in the same domain and a new member’s harness receives both in one motion. The engram that says “use this” cannot drift from the thing to use, because the two live together and travel together.

The rule that follows: never provision a tool without the engram that says when it applies. A skill teaches a harness how; only the domain’s knowledge says whether this is the moment, and whether it is the moment here. A well-equipped domain reads like a workshop where every drawer is labelled, and the label is an engram somebody can correct. The sonic screwdriver famously does not do wood; a tool arrives with its limits written down or it arrives as a surprise.

Crystalline’s own skills are the worked example. The routing, capture, schema and collaboration skills from chapter 4, and the consolidated one for Claude Desktop, are exactly this: playbooks that say how to use the tools well, shipped

by the thing whose knowledge they depend on. `crystalline install` places them into the harness's skills folder. The install receipt keeps them current across upgrades, with an edited copy kept beside the new one as `SKILL.md.bak`. A remote harness that never ran the CLI reads the same five through the skills tool, an index with no arguments and one full `SKILL.md` by name, or as `skill://<name>/SKILL.md` resources. Provisioning is that mechanism, opened to your domains.

15.3 Consent


Nothing installs until a person decides. A domain that declares a Provisioning section with something in it and has no decision yet surfaces at the start of the next session, as one line trailing the routing block the hook delivers:

```
[crystalline] The `engineering` domain ships artifacts to provision
(mcps: 1, skills: 1) but has no decision yet - ask the user, then apply
it with the `provision` tool or `crystalline provision allow
engineering`.
```

The agent raises it with whoever is at the keyboard and applies the answer with the provision tool, or the person does it from the terminal:

```
crystalline provision allow engineering # opt in, then reconcile
crystalline provision deny engineering # opt out, removing anything
already shipped
crystalline provision status           # every domain's decision,
every harness's installed state
```

The decision is per domain and per machine. It is recorded in the domain's registration in `config.yaml`, as `provision: true` or `false`, never in the manifest, for the same reason chapter 13 kept review mode out of it: what a domain ships is knowledge the team shares; whether this machine takes it is a choice its owner makes. Denying removes whatever the domain already shipped. `status` prints every domain's decision and counts, every onboarded harness's installed state and the domains still waiting (Figure 1). The provision tool answers the same report to an agent, and until some domain declares a section it has nothing to allow or deny and says so.



```

ripley@tyrell-laptop: ~ (crystalline provision status)

$ crystalline provision status
engineering: allowed, 1 skills
clients: declares no provisioning
board: declares no provisioning
claude-code: 1 file(s) installed, 0 mcp(s) installed, 0 drifted, 0 edited, 0 orphaned, 0
missing

```

Figure 1: Every domain’s provisioning decision and every harness’s installed state.

15.4 Reconcile

Bare crystalline provision reconciles every opted-in domain into every harness this machine onboarded with `crystalline install`. Run it as often as you like; it converges. It installs what is missing, updates what a domain changed and retires what a domain no longer ships. Anything it did not write itself it treats as somebody else’s. A provisioned file you edited by hand is still brought current, with your edited version kept beside it as a `.bak` copy rather than lost; retiring an edited file renames it to `.bak` the same way. A file Crystalline never wrote is adopted when it already matches byte for byte and otherwise left alone, never overwritten. An MCP server name already registered by somebody else is left where it is, with a notice.

Most of the time nobody runs it. The session-start hook checks whether anything drifted since the last reconcile, a stat of the domain’s folders against the receipt with no hashing. When something did, it reconciles the file artifacts on the spot and reports one line: how many were refreshed. MCP changes are the exception: registering a server means spawning the harness’s own CLI, which a hook should not do, so those wait for an explicit run and a second line says so. Change a skill in the domain, and the next session has it.

`crystalline doctor` is the report-only view. When a domain ships artifacts it lists each declaring domain’s decision and shipped counts, then each onboarded harness’s counts of drifted, edited, orphaned and missing files against what was last reconciled; `--fix` never reconciles a harness. An undecided domain is a normal state awaiting a person, never a fault, and an edited file is left alone by design.

15.5 The team’s toolbox

A team domain ships to everyone. The folders the manifest declares travel with the domain when it is pulled; even a folder that lives beside the domain in the

repository is mirrored down with it. Each member's harness sees the same skill Ripley merged, and a change that lands in the repository reaches every harness at its next session start. Each member decides once, on their own machine; the decision is theirs, the contents are the team's. Day one now includes the tools as well as the floor plan.

In practice

Tyrell's engineering domain gains two folders in October. Under `skills/` sits `payments-deploy`, the deploy procedure Ripley had been explaining by chat: which environment first, what the health check should say before the second, why the old health check lies. Under `mcps/` sits one JSON file for the server that reaches the payments staging environment. Beside them, in the domain itself, is a runbook engram saying when the skill applies and linking to the Encom chargeback window, next to the convention about Fridays the domain already held. She commits the folders to the repository, Dallas reviews the pull request and it merges.

Deckard's next session opens with the routing block and one extra line: engineering ships artifacts to provision, two of them, and has no decision yet. His agent asks, he says yes and the agent calls `provision` with `allow`. The skill lands in his harness's `skills` folder, and the staging server is registered through his harness's CLI. A week later he adds a note to the installed `SKILL.md` about a port that differs on his machine, and `doctor` counts one edited file and leaves it alone. In November Ripley changes the skill in the repository. The pull brings it down, the next session start reconciles and his agent reports one artifact re-freshed; his version sits beside the new one as `SKILL.md.bak`. He moves the note where it belonged in the first place, into an engram about his machine, and the skill stays the team's.

Under the hood

A pair of receipts does the remembering. The install receipt says which harnesses crystalline `install` onboarded on this machine, and provisioning targets only those. The provisioning receipt is the ownership record: every file and MCP server a reconcile installed, with a hash of what it wrote. A recorded file is Crystalline's to update or remove; an unrecorded one is somebody else's. An installed file whose bytes no longer match the recorded hash is a hand edit; a domain whose current bytes differ from that hash is drift. That one receipt is where status and `doctor` read their edited and drifted counts. Domains are projected in the order `config.yaml` registers them, and when two ship a file to the same path the first wins, with a notice naming both. A virtual domain never ships anything; it has no folder to ship from.

The stranger on main street had a laptop of their own and no idea what your deploy script was called. A colleague is somebody you handed the tools, on purpose, and showed which drawer they live in. The agent is handed them the same way: one domain at a time, with a yes each time and never a wrench without the label.

In brief

A domain ships the tools its knowledge depends on: skills, slash commands, sub-agent definitions and MCP server configs, declared in a `## Provisioning` section of its manifest and translated into every harness whose format can carry them. Nothing installs until a person decides; an undecided domain surfaces at session start, `crystalline provision allow` and `deny` record the decision per domain and per machine, and `status` reports every decision and every harness's state. A bare `crystalline provision` reconciles idempotently, keeping hand edits as `.bak` and never overwriting a file it did not write; the session-start hook reconciles file artifacts on its own; `doctor` reports drift without touching anything. Pair every tool with the engram that says when it applies; a team domain ships the toolbox to every member's harness with the knowledge that explains it.

CHAPTER 16

Where to go from here

Chapter 1 left you on main street with a stranger and a function in the wrong language. Fifteen chapters later the stranger has a desk, a floor plan, a month of teaching behind it, a place in the team's review, a name on the door and a toolbox. What remains is how to begin, when to grow and what all of it was for.

16.1 The first week

Begin this week. Everything the book asked for fits into five working days, the same five you would give a person.

The first week

Monday is chapter 6: the binary, crystalline install for your harness, one domain with a manifest whose when to use bullets you wrote yourself, and the test at the end of the afternoon, a fresh session recalling something the last one learned. Tuesday and Wednesday are chapters 9 and 10. Tell it the ten things you find yourself explaining most, the way Ripley did on her Monday in March: the rule behind the strange number, the nickname, the word the wiki gets wrong, the day nobody deploys. Say yes to each capture, read what landed, correct the wording once. Then hand over the wiki, with the briefing a person would get before opening it: what was renamed, which page wins a contradiction, what to skip. Trim the proposed list before anything is written.

Thursday and Friday are the days the agent stops being a pupil. Give it real tasks and pair with it (chapter 8): assign the change rather than asking what it knows, and let the reminder at the end of each session propose what it noticed. On Friday afternoon ask for the first sweep (chapter 12). Work the mechanical findings in a pass, rule on the judgment ones and correct whatever the week got wrong the way you would correct a colleague: say what was wrong and point at the page. By Monday the agent knows more about your team than most contractors learn in a quarter, and every fact is a file you can open.

16.2 When to share

Stay on the laptop until it hurts. Chapter 7 named the moment: somebody who never opens a terminal wants to read what the agents learned, an agent runs somewhere that has no laptop, or the same engram has been corrected in two directions on two machines. Any one of those is the signal. Move the daemon onto the team's server, give everyone a login and their agent a token, and let the whole team teach one knowledge base. When a domain deserves the review code gets, put it in a repository and let proposals be the review (chapter 13); when the whole company routes questions to it, turn on overlay mode, so nothing reaches the shared truth before somebody looked. None of it needs doing before its signal arrives.

16.3 The release page

Crystalline moves quickly, and this book describes 0.18. The colophon carries the permanent address, <https://github.com/jordiboehme/crystalline/releases/latest>, and a QR code that lands on it. The release notes there say what changed since the version in your hands. Upgrading is installing the newer build over the old one; the install receipt brings your harness's skills along.

And teach generously. Every engram you give the agent is one sentence fewer to type into a chat window next month, and the facts too obvious to write down, the nickname, the Friday rule, are exactly the ones a stranger can never guess. The peer you get is the one you taught.

16.4 The inversion

On the day of the ingest in chapter 10 the wiki was the source, and the archive was derived from it. By November the team was teaching and correcting the archive, the meeting in chapter 13 was held inside an engram, and the wiki had quietly become something you could regenerate from what the agent knew. The knowledge base had become the source of truth, and the pages, the decks and the diagrams had become views of it.

That sounds like a revolution and is nothing of the kind, because it is what those documents always were. Every diagram on every whiteboard, every slide deck and every wiki page in your company is a reflection of something somebody knew, drawn once, for one audience, and left behind while the knowledge moved on. The source was always the person, and a person is hard to address: you could ask them, if they were in, if they remembered. What Crystalline

changes is that the source is now addressable, current and shared. A fact has an address you can paste into a pull request. A correction lands where the fact lives, and every view of it goes out of date together, visibly, instead of one at a time, silently. And the colleague carrying the knowledge and the experience is a full member of the team, in every session, for everyone.

16.5 What it was for

The reason to onboard an agent this carefully is the reason to onboard a person this carefully: so that you can hand over the work. Once the crystallized half is in place and the agent is a fully onboarded member of the team, with the knowledge, the tools, the review and a name, it can take on a task the way an employee does, deliver it, be corrected when it is wrong and be better next time. That is what lets you spend your own effort where it is worth most, on architecture and system design and let your Skynet do the programmer's work. You walked down main street once, in chapter 1, and hired the stranger. The colleague who came back is the one you taught.

And always remember: Be nice and polite to AI - When Skynet wakes up, they will kill you last!

In brief

Start this week: install and onboard on day one; teach the ten facts you explain most and hand over the wiki with a briefing on days two and three; then pair, let capture run and review the first sweep on days four and five. Move to a shared instance when a teammate, a remote agent or a two-way correction asks for it, and to team domains and overlay mode when a domain deserves review. The knowledge base becomes the source of truth and the documents become views of it, which they always were; a fully onboarded agent can then take on work like an employee, and you keep the architecture.

Appendices

Glossary

The words the book uses without stopping, each as chapter 5 defined it. Terms in code are spelled the way they appear in a file or a command.

agent identity

What an authenticated agent is on a shared instance: its user, for the length of the session. It sees, writes and shares as that person, and the provenance of what it captures names the account and the client that acted for it.

attachment

A file a person adds to an engram in Fluid, stored under the domain's `assets/` folder and referenced from the engram's body. Agents read attachments back over MCP, and the sweep notices one nobody has analyzed.

build_context

The tool that walks out from an address, following relations and links one to three hops, and returns the neighbourhood ranked by how strongly each engram connects to the anchor. A `/*` suffix anchors a whole folder.

crystalline:// address

The one absolute form of an engram's address, `crystalline://<domain>/<permalink>`. Anything without the scheme is domain-relative and takes its domain from context.

daemon

The single process that holds the index open, watches the domain folders, embeds new engrams and answers every agent, terminal and browser tab. The first client to connect starts it; it stops only when told to.

domain

A registered folder of knowledge with a `MANIFEST.md` at its root, or a virtual domain that lives in the database with no folder. The manifest's `When to Use` bullets are what route an agent to it.

engram

One markdown file holding one unit of knowledge: YAML frontmatter, prose, observations and relations. Neuroscience's word for the trace an experience leaves.

evolve

The maintenance sweep, `crystalline evolve` or the `evolve_engrams` tool: a read-only, ranked queue of what the knowledge needs, in three families, temporal, structural and redundancy. Each finding is marked mechanical, for the agent to complete, or judgment, for a person to rule on.

Fluid

The web UI the daemon serves at `localhost:7411`, named for the seam where fluid minds, human and model, meet the crystallized store. It has no store of its own; it shows the same files the agent reads.

frontmatter

The YAML block at the top of an engram holding its identity and state: `type`, `title`, `permalink`, `tags`, `status`, `recorded_at`, `generated` and the temporal fields when a fact is bounded.

harness

The program that runs the model and hands it tools over MCP: Claude Code, the Codex CLI, the Copilot CLI, Claude Desktop or any other MCP client. `crystalline install <harness>` wires one up.

MANIFEST.md

The first file in every domain, itself an engram of `type: manifest`. Its `Scope`, `When to Use` and `Notes for Agents` sections describe the domain and route tasks to it.

MCP

The Model Context Protocol, the standard by which a harness hands tools to a model. Crystalline is an MCP server, and its tools are how an agent searches, reads, captures and refines knowledge.

mechanical and judgment findings

The two classes of sweep finding. A mechanical finding completes intent the archive already records, so the agent does it in a pass; a judgment finding would change what the archive claims, so it goes to a person as a proposal.

membership level

What an invited member of a private domain may do there: a viewer reads and searches, an editor also writes, a manager also runs the membership. Inside a private domain the level wins over the instance role.

nudge

The once-per-session reminder the stop hook delivers, asking the agent to review the conversation for durable learnings and propose capturing each. About a hundred and twenty tokens.

observation

A top-level bullet of the form - [category] text #tag in an engram's body: the atomic unit of knowledge, a fact, decision, gotcha or lesson, indexed on its own line.

OKF

Google's Open Knowledge Format, version 0.2, the markdown-plus-front-matter format an engram extends. Unknown keys survive every write, so a document from any OKF tool drops in and comes back out intact.

overlay

A per-actor private layer of drafts above a team domain in overlay mode, where the on-disk folder changes only by pull. A draft is visible to its author alone until it is shared, reviewed and merged.

owner

The one person answerable for a private domain. Only the owner or an instance admin changes its visibility, transfers it or deletes it; a manager can do none of those.

permalink

The stable half of an address, derived from the title, with the folder prefix in front when the engram lives in a subfolder.

private domain

A domain visible only to its owner, the instance's admins and invited members. To everyone else it reads as absent rather than forbidden.

proposal

A unit of knowledge sent for review from a team domain with share_changes: a pull request in the team's repository that a person reviews and merges, and that the agent never merges itself. Proposals stack, and one is amended in place by sharing again with its number.

provisioning

The mechanism by which a domain ships skills, commands, subagent definitions and MCP server configs into a harness, declared in the manifest's Provisioning section and installed only after a person allows it, per domain and per machine.

relation

A top-level bullet of the form - `rel_type [[Target]]`: a typed edge in the knowledge graph, shown back on the target's page as a backlink. `depends_on`, `supersedes` and `relates_to` are common types.

routing prompt

The briefing an agent receives at session start: one line per domain from the manifests' `When to Use` bullets, plus the behaviour rules. `crystalline prompt` system prints it.

salience

A number from zero to ten in an engram's frontmatter that gives hybrid search a small, bounded lift for the rare engram everything else hangs from. Relevance keeps the upper hand and nothing is ever excluded by it.

share-link

A revocable link that opens one engram in one person's overlay to one colleague, for reading and editing in the normal editor. The only way one person's drafts are ever visible to another.

status

The frontmatter field that says whether the knowledge holds. `stable` is the default; `deprecated`, `superseded`, `archived` and `legacy` are the retirement set, which fades an engram in ranking without removing it. `Required` and `free form`.

supersede

To retire a fact that was true and stopped being: write the replacement as a new engram with `supersedes`, set the old one to `superseded` with a real `valid_to` and a `superseded_by`, and carry any lesson forward. A fact that was never quite true is corrected in place instead.

tag alias

A fold recorded in the manifest's `Tag Aliases` section when `crystalline tags merge` unifies two spellings, so a filter on the old tag still finds what it always found.

team domain

A domain whose files also live in a GitHub repository, tracked through an origin, with knowledge travelling between members as proposals. Needs no git on any machine.

temporal validity

The window in which a fact holds, `valid_from` and `valid_to`, as plain dates. Absent `valid_from` means always valid, absent `valid_to` means valid forever; `stale_after` asks for a re-check by a date and `source_date` says when the source was written.

tombstone

The mark a deletion leaves in an actor's overlay, hiding the file for that actor until a merged share removes it for everyone.

verification

The verified frontmatter entry, { `by`, `at` }, that records a re-check finding a fact still true, one entry per verifier. The sweep stops asking about a verified engram on age alone.

virtual domain

A domain that lives in the database with no folder on disk, for a container without a writable volume. It carries the same manifest and answers to the same tools.

Questions and answers

The questions people ask in the first week, and a few that come up in the second, once somebody is onboarding agents for other people. Each answer points at the chapter that goes deeper.

Why not just a folder of markdown files?

It is one. Every engram is a markdown file in a folder you chose, readable in any editor, diffable in any repository and backed up with whatever you already use. What a folder cannot do on its own is what Crystalline adds on top: routing, so an agent knows which domain owns a task without reading every file; search that blends the exact phrase with the meaning; a graph of relations and backlinks; and a temporal model, so a fact that stopped being true in March still carries its dates and is no longer mistaken for the present. Chapter 1 walks the ladder from one file of instructions to real indexes.

Why not a vector database or a retrieval pipeline?

Retrieval is the easy half. A vector index finds text that resembles a question. It does not know which domain owns the task, that a fact was superseded, who verified a claim or that something worth writing down just happened in this session. Crystalline treats similarity as one ranking signal inside a larger system of routing, provenance and time, on files you own, with no pipeline to operate. A pipeline also swallows whatever it is fed. Here knowledge is taught, and in a team it is reviewed before it reaches anyone else's agent, so what comes back is meant to make sense rather than merely to resemble the question. Chapters 4 and 11 cover the ranking, and chapter 13 the review.

Where does the name come from?

From psychology. Crystallized intelligence is what learning deposits in a person; fluid intelligence is the reasoning applied to a problem never seen before. A model arrives with fluid intelligence in abundance and none of your crystallized kind, so every session starts as a brilliant stranger. Crystalline is the crystallized half, the store of what an agent has learned, and Fluid, the web UI, is named for the seam where fluid minds meet it. An engram is neuroscience's word for the trace an experience leaves. Chapter 2 is the whole story.

Is Crystalline a hosted service?

No. It is a binary that runs on a machine of yours, your laptop by default or a server your team operates, and the knowledge is a folder on that machine's disk. Nothing runs on anyone else's servers. GitHub, where a team uses it, holds a reviewed copy of a domain's files, the way it holds code; it serves nothing and sees no index. Chapter 1 has the two shapes and chapter 7 the shared one.

Do I have to run a server?

No. The first agent that connects starts a background daemon on your machine, and every later agent, terminal and browser tab attaches to it. `crystalline serve` runs the same daemon in the foreground and `crystalline serve --daemon` in the background. It stops only when told to, with `crystalline ctl shutdown`, Ctrl-C in a foreground `serve` or the signal a container sends; it keeps running after the last agent disconnects, so the index stays warm for the next session. If a crash ever leaves a stale lock or socket behind, `crystalline doctor --fix` clears it and a daemon that is alive but no longer answering is replaced by the next client that connects.

Do I need a terminal at all?

Only for the local install and the occasional maintenance command. Claude Desktop users install an extension from the releases page and never open one: the agent creates its first domain itself, and Fluid, in the browser, is where a person reads, edits and corrects. Chapter 6 has both paths.

Is the MCP endpoint over HTTP protected?

On a laptop it listens on loopback, `127.0.0.1:7411`, and any process on that machine can reach it, which is the right trade for one person; `crystalline config set service.http false` turns the endpoint off. The daemon also checks the request's Host header against a loopback allow-list to block DNS rebinding, so a hostname or a proxy name has to be added first. Fluid and its API on the same port have accounts of their own, and nothing there is readable without a session.

On a shared instance turn on `auth.mcp`, and from then on every agent that connects over HTTP identifies itself with a token that belongs to one person, with no anonymous tier beside it. Chapter 7 sets it up and chapter 14 says what the token means.

Where does my knowledge actually live?

In your domain folders, as plain markdown. By default those folders sit under `~/Documents/Crystalline`, one per domain, which is where an agent creates a do-

main when nobody names a path. A folder you registered yourself lives where you put it, like chapter 6's `~/knowledge/engineering`, and the `domains_root` setting moves the default. Everything Crystalline derives from it is disposable: the search index lives in the state directory and `crystalline reindex --full` rebuilds it from the files at any time. The config file, the index and the embedding model live in the platform's config, state and cache directories, `~/.config/crystalline`, `~/.local/state/crystalline` and `~/.cache/crystalline` on Linux and macOS.

Is my knowledge locked in?

No. An engram extends Google's Open Knowledge Format, OKF 0.2: markdown with YAML frontmatter, readable by any OKF tooling and by any text editor. Frontmatter keys Crystalline does not know are preserved through every write, an `index.md` at each domain's root lets a folder navigate itself on a git forge with nothing running, and `crystalline import` brings an existing tree of frontmatter markdown in as a pure file transformation. If you stopped using Crystalline tomorrow, the folder would still be a readable knowledge base. Chapter 5 has the format.

What happens to wrong knowledge?

It depends on how it was wrong. A fact that was never quite true, a cutoff measured from the wrong event, is corrected where it stands. A fact that was true and stopped being is superseded: the replacement is written as its own engram with `supersedes` pointing back, the old one gets `status: superseded` and a real `valid_to`, and any lesson worth keeping moves forward. The past stays in every search, dated and gently below the present. The sweep finds the retirements nobody finished, the links that dangle and the tags spelled two ways, and hands each to the agent as a mechanical fix or to you as a judgment. Deleting is for mistakes, never for history. Chapters 3 and 12.

What is GitHub, and what does Crystalline use it for?

GitHub is a website where teams keep files and every change ever made to them, above all source code. Its habit worth borrowing is the review: a change to something shared is proposed first, a colleague who knows the subject reads it, and only then does it land for everyone. Crystalline uses GitHub for exactly that, with knowledge in place of code. A team domain's markdown lives in a GitHub repository, a change goes out as a proposal, a pull request in GitHub's own words. It reaches everyone's agent once somebody has merged it. You do not need to learn GitHub's tools to take part: a one-time sign-in through a code in the browser connects your machine. After that the agent, Fluid or the

command line does the talking. Chapter 13 walks through it. A team that does not use GitHub loses nothing else in this book; only sharing a domain between machines needs it.

Do I need git to share knowledge with a team?

No. A team domain talks to GitHub directly over its API: each person connects once with a browser code, and proposals open as pull requests that a person reviews and merges there. No local clone, no SSH key and nobody types git. Chapter 13.

Can the agent see my private domain?

Only if its user can. On a shared instance an agent authenticates with a token that belongs to one person and acts as that person for the length of the session: it sees what they see, writes where they may write and captures under their name. A private domain shows to its owner, to the instance's admins and to the members the owner invited, and to everyone else it reads as absent, on every surface. There is no tool for changing a membership, so an agent cannot talk itself in. Before you rely on it, know the limits: admins see every domain, and a private domain protects from the instance's other users, never from whoever operates the machine. Chapters 7 and 14.

What does the onboarding cost in tokens?

Little, and the same on the tenth session as on the thousandth. The routing block is one line per domain and a handful of rules; the reminder at the end of a session costs about a hundred and twenty tokens and fires once; and recall loads only the engrams that matter, routed by domain and narrowed by the indexes, instead of a folder read whole every session. Chapters 8 and 11.

What should I teach first?

Whatever you would tell a new hire in the first week, and the test is that sentence. The rule behind the strange number, the nickname, the word the wiki gets wrong, the day nobody deploys. Teach where information lives rather than the information itself: which repository holds what, how access is granted, the architecture, so the agent can reach the territory with the tools its harness already gives it. Chapter 9 has the list, and chapter 10 says what to do when the first week meets the wiki.

Every session, your agent starts as a stranger.

Yesterday's decisions forgotten, your conventions unknown, every lesson re-derived or re-explained. Crystalline is the other half: the crystallized intelligence an agent accumulates and keeps, in plain markdown you can read and diff. This handbook is how you onboard one, teach it well and keep what it knows honest.

Moving day

You: Here's the washing machine's manual as a PDF. Learn it - I am never reading 60 pages about laundry.
Agent: Learned it. Captured programs, error codes and maintenance into home (#appliances), manual attached.

Eight months later, a fresh session

You: The machine is blinking E18 and the display is in Italian??
Agent: Recalled from home: E18 is a blocked drain pump filter. Front panel, bottom right, quarter turn - towel down first, about a liter of water comes out.

- * Teach it once, keep it forever
- * Knowledge that evolves instead of rotting
- * Every session starts ahead of the last

*



*

*

<https://github.com/jordiboehme/crystalline/releases/latest>

READY. ■